

# Evolution in Model-Driven Engineering

Louis M. Rose

November 3, 2009



# Contents

<b>1</b>	<b>Introduction</b>	<b>5</b>
1.1	Model-Driven Engineering . . . . .	5
1.2	Software Evolution . . . . .	5
1.3	Research Aim . . . . .	5
1.4	Research Method . . . . .	5
<b>2</b>	<b>Background</b>	<b>7</b>
2.1	Related Areas . . . . .	7
2.2	Categories of Evolution in MDE . . . . .	7
<b>3</b>	<b>Literature Review</b>	<b>9</b>
<b>4</b>	<b>Analysis</b>	<b>11</b>
4.1	Locating Data . . . . .	11
4.1.1	Requirements . . . . .	11
4.1.2	Project Selection . . . . .	12
4.1.3	Other examples . . . . .	14
4.1.4	Summary . . . . .	17
4.2	Analysing Existing Techniques . . . . .	17
4.2.1	Modelling Framework Characteristics Relevant to Co-Evolution	17
4.2.2	Categorisation of Co-Evolution Techniques . . . . .	18
4.3	Requirements Identification . . . . .	20
<b>5</b>	<b>Implementation</b>	<b>19</b>
5.1	Metamodel-Independent Syntax . . . . .	19
5.2	Human-Usable Textual Notation . . . . .	19
5.3	DSL for Migration Strategies . . . . .	19
<b>6</b>	<b>Evaluation</b>	<b>21</b>
6.1	Case Study . . . . .	21
6.2	Publications . . . . .	21
6.3	Delivery through Eclipse . . . . .	21
6.4	Limitations . . . . .	21
<b>7</b>	<b>Conclusion</b>	<b>23</b>



# Chapter 4

## Analysis

### 4.1 Locating Data

Two categories of evolutionary change encountered in model-driven engineering were identified in Chapter 2. Existing MDE projects were examined for examples of evolution, which would be used to provide requirements for developing structures and processes for evolutionary changes in the context of MDE. In this section, the requirements used to select example data are described, along with candidate and selected MDE projects. The section concludes with a discussion of further examples, which were obtained from collaborative research and from a related software engineering domain.

#### 4.1.1 Requirements

In Chapter 2, two categories of evolutionary change were identified: co-evolution and synchronisation. In this section, requirements are partitioned into three types: those necessary for studying each of the two categories of evolutionary change, and common requirements (applicable to both categories of evolutionary change).

##### Common requirements

Every candidate project needs to use MDE. Specifically, both metamodeling and model transformation must be used (requirement R1). In addition, each candidate project needs to provide historical information to trace the evolution of development artefacts (R2). For example, several versions of the project are needed perhaps in a source code management system. Finally, a candidate project needs to have undergone a number of significant changes<sup>1</sup> (R3).

##### Co-evolution requirements

A candidate project for the study of co-evolution needs to define a metamodel and some changes to that metamodel (R4). In the projects considered, the metamodel changes took the form of either another version of the metamodel, or a history (which recorded each of the steps used to produce the adapted

---

<sup>1</sup>This is deliberately vague. Further details are given in Section 4.1.2.

Table 4.1: Candidates for study of evolution in existing MDE projects

Name	Requirements									
	Common			Co-evolution			Synchronisation			
	R1	R2	R3	R4	R5	O1	R6	R7	R8	O2
GSN	x			x						
OMG	x			x			x			
Zoos	x	x		x						
MDT	x	x		x		x				
MODELPLEX	x	x	x	x		x	x	x		
FPTC	x	x	x	x	x					
xText	x	x	x	x	x	x	x	x		x
GMF	x	x	x	x	x	x	x	x		x

metamodel). A candidate project also needs to provide example instances of models before and after each migration activity (R5).

Ideally, a candidate project should include more than one consecutive meta-model adaptation, so as to represent the way in which the same development artefacts continue to evolve over time (optional requirement O1).

#### Synchronisation requirements

A candidate project for the study of synchronisation needs to define a model-to-model transformation (R6). Furthermore, a candidate project has to include many examples of source and target models for that transformation (R7). Crucially, a candidate project needs to provide many examples of the kinds of change (to either source or target model) that cause inconsistency between the models (R8).

Ideally, a candidate project should also include transformation chains (more than one model-to-model transformation, executed sequentially) (O2). Chains of transformations are prescribed by the MDA guidelines [4].

#### 4.1.2 Project Selection

Eight candidates were considered for the study. Table 4.1 shows which of the requirements are fulfilled by each of the candidates. Each candidate is now discussed in turn.

##### GSN

Georgios Despotou and Tim Kelly, members of this department’s High Integrity Systems Engineering group, are constructing a metamodel for Goal Structuring Notation (GSN). The metamodel has been developed incrementally. There is no accurate and detailed version history for the GSN metamodel (requirement R2). **Suitability for study:** Unsuitable.

##### OMG

The Object Management Group (OMG) [8] oversees the development of model-driven technologies. The Vice President and Technical Director of OMG, An-

drew Watson, references the development of two MDE projects in [13]. Correspondence with Watson ascertained that source code is available for one of the projects, but there is no version history. **Suitability for study:** Unsuitable.

### Zoos

A zoo is a collection of metamodels, authored in a common metamodeling language. I considered two zoos, but neither contained any significant external metamodel changes. Those changes that were made involved only renaming of meta-classes (trivial to migrate) or additive changes (which do not affect consistency, and therefore require no migration). **Suitability for study:** Unsuitable.

### MDT

The Eclipse Model Development Tools (MDT) [1] provides implementations of industry-standard metamodels, such as UML2 [7] and OCL [6]. Like the metamodel zoos, the version history for the MDT metamodels contained no significant changes. **Suitability for study:** Unsuitable.

### MODELPLEX

Jendrik Johannes, a research assistant at TU Dresden, has made available work from the European project, MODELPLEX. Johannes's work involves transforming UML models to Tool Independent Performance Models (TIPM) for simulation. Although the TIPM metamodel and the UML-to-TIPM transformation have been changed significantly, no significant changes have been made to the models. **Suitability for study:** Unsuitable.

### FPTC

Failure Propagation and Transformation Calculus (FPTC), developed by Malcolm Wallace in this department, provides a means for reasoning about the failure behaviour of complex systems. Before starting my doctorate, I worked with Richard Paige to develop an implementation of FPTC in Eclipse. The implementation includes an FPTC metamodel. Recent work with Philippa Conmy, a research assistant in the department, has identified a significant flaw in the implementation, leading to changes to the metamodel. These changes caused existing FPTC models to become inconsistent with the metamodel. Conmy has made available copies of FPTC models from before and after the changes. **Suitability for study:** Suitable for studying co-evolution. Unsuitable for studying synchronisation.

### xText

xText is an openArchitectureWare (oAW) [9] tool for generating parsers, metamodels and editors for performing text-to-model transformation. Internally, xText defines a metamodel, which has been changed significantly over the last two years. In several cases, changes have caused inconsistency with existing models. xText provides examples of use, which have been updated alongside the metamodel. **Suitability for study:** Suitable for studying co-evolution. Unsuitable for studying synchronisation.

## GMF

The Graphical Modelling Framework (GMF) [3] allows the definition of graphical concrete syntax for metamodels that have been defined in EMF. GMF prescribes a model-driven approach: Users of GMF define concrete syntax as a model, which is used to generate a graphical editor. In fact, five models are used together to define a single editor using GMF.

GMF defines the metamodels for graphical, tooling and mapping definition models; and for generator models. The metamodels have changed considerably during the development of GMF. Some changes have caused inconsistency with GMF models. Presently, migration is encoded in Java. Gronback has stated<sup>2</sup> that the migration code is being ported to QVT (a model-to-model transformation language) as the Java code is difficult to maintain.

GMF fulfils almost all of the requirements for the study. A large amount of the co-evolution data is available, including migration strategies. The GMF source code repository does not contain examples of the kinds of change that cause inconsistency between the models. However, GMF has a large number of users, and it may be possible to gather this information elsewhere. **Suitability for study:** Suitable for studying both categories of evolutionary change.

## Summary of selection

The FPTC and xText projects were selected for a study of co-evolution. No appropriate projects were located for a study of synchronisation. The GMF project will not be studied immediately, but reserved for evaluation (Chapter 6).

### 4.1.3 Other examples

Because only a small number of MDE projects fulfilled all of the requirements, additional data was collected from alternative sources. Firstly, examples were sought from a related domains, object-oriented systems. Secondly, examples were discovered during collaboration with colleagues on two projects, both of which involved developing a system using MDE.

## Examples of evolution from object-oriented systems

In object-oriented programming, software is constructed by developing groups of related objects. Every object is an instance of (at least) one class. A class is a description of characteristics, which are shared by each of the class's instances (objects). A similar relationship exists between models and metamodels: metamodels comprises meta-classes, which describe the characteristics shared by each of the meta-class's instances (elements of a model). Together, model elements are used to describe one perspective (model) of a system. This similarity between object-oriented programming and metamodeling implied that the evolution of object-oriented systems may be similar to evolution occurring in MDE.

*Refactoring* is the process of improving the structure of existing code while maintaining its external behaviour. When used as a noun, a refactoring is one such improvement. As discussed in Chapter 3, refactoring of object-oriented

---

<sup>2</sup>Private communication, 2008.



systems has been widely studied, perhaps most notably in [2], which provides a catalogue of refactorings for object-oriented systems. For each refactoring, Fowler gives advice and instructions for its application.

To explore their relevance to MDE, the refactorings described in [2] were applied to metamodels. Some were found to be relevant to metamodels, and could potentially occur during MDE. Many were found to be irrelevant, belonging to one of the following three categories:

1. **Operational refactorings** focus on restructuring behaviour (method bodies). Most modern MDE modelling frameworks do not support the specification of behaviour in models.
2. **Navigational refactorings** convert, for example, between bi-directional and uni-directional associations. These changes are non-breaking in EMF, which automatically provides values for the inverse of a reference when required.
3. **Domain-specific refactorings** manage issues specific to object-oriented programming, such as casting, defensive return values, and assertions. These issues are not relevant to metamodeling.

The object-oriented refactorings that can be applied to metamodels provide examples of metamodel evolution. When applied, some of these refactorings potentially cause inconsistency between a metamodel and its models. By using Fowler's description of each refactoring, a migration strategy for updating (co-evolving) inconsistent models was deduced. An example of this process is now presented.

Figure 4.1 illustrates a refactoring that changes a reference object to a value object. Value objects are immutable, and cannot be shared (i.e. any two objects cannot refer to the same value object). By contrast, reference objects are mutable, and can be shared. Figure 4.1 indicates that applying the refactoring restricts the multiplicity of the association (on the Order end) to 1 (implied by the composition); prior to the refactoring the multiplicity is many-valued.




Figure 4.1: Refactoring a reference to a value. Taken from [2][pg183].

Before applying the refactoring, each customer may be associated with more than one order. After the refactoring, each customer should be associated with only one order. Fowler indicates that every customer associated with more than one order should be duplicated, such that one customer object exists for each

order. Therefore, the migration strategy in Listing 4.1 is deduced. Using this process, migration strategies were deduced for each of the refactorings that were applicable to metamodeling, and caused inconsistencies between a metamodel and its models.

Listing 4.1: Migration strategy for the refactoring in pseudo code.

```

for every customer, c
  for every order, o, associated with c
    create a new customer, d
    copy the values of c's attributes into d
  next o

  delete c
next c

```

The examples of metamodel evolution based on Fowler's refactorings provided additional data for deriving research requirements. Some parts of the metamodel evolutions from existing MDE projects were later found to be equivalent to Fowler's refactorings, which, to some extent, validates the above claim that evolution from object-oriented systems can be used to reason about metamodel evolution.

On the other hand, object-oriented refactorings are used to improve the maintainability of existing systems. In other words, they represent only one of the three reasons for evolutionary change defined by [12]. The two other types of changes are equally relevant for deriving research requirements, and so object-oriented refactorings alone are not sufficient for reasoning about metamodel evolution.

### Research collaborations

As well as the example data located from object-oriented system, collaboration on projects using MDE with two colleagues provided several examples of evolution. A prototypical metamodel to standardise the way in which process-oriented programs are modelled was produced with Adam Sampson, a research assistant at the University of Kent, and an investigation of the feasibility of implementing a tool for generating story-worlds for interactive narratives was conducted with Heather Barber, a postdoctorate student in the department.

In both cases, a metamodel was constructed for describing concepts in the domain. The metamodels were developed incrementally and changed over time. The collaborations with Sampson and Barber did not involve constructing model-to-model transformations, but did provide data suitable for a study of co-evolution.

Majority of the changes made in both of these projects related to changing requirements. In each iteration, existing requirements were refined and new requirements discovered. Neither project required changes to support architectural restructuring. In addition, the work undertaken with Sampson included some changes to adapt the system for use with a different technology than originally anticipated. That is to say, the changes observed represented two of the three reasons for evolutionary change defined by [12].

#### 4.1.4 Summary

To summarise, in this section the example data used to analyse existing structures and processes for managing evolution in the context of MDE was discussed. Example data was sought from existing MDE projects, and also from a related domain and research collaboration. Eight existing MDE projects were located, three of which satisfied the requirements for a study of co-evolutionary changes in the context of model-driven engineering. One of the three projects, GMF, was reserved as a case study and is used for evaluation in Chapter 6. Refactorings of object-oriented programming supplemented the data available from the existing MDE projects. Collaboration with Sampson and Barber yielded further examples of co-evolution.

Due to the lack of examples of model synchronisation, this thesis now focuses on model and metamodel co-evolution.

## 4.2 Analysing Existing Techniques

The examples of co-evolution identified in the previous section were analysed to discover and compare existing techniques for managing co-evolution. This section discusses the results of analysing the examples; namely a deeper understanding of modelling framework characteristics that affect the management of co-evolution and a categorisation of existing techniques for managing co-evolution. These results were published in [11, 10].

### 4.2.1 Modelling Framework Characteristics Relevant to Co-Evolution

Analysis of the co-evolution examples identified in the previous section highlighted characteristics of modern MDE modelling development environments that impact on the way in which co-evolution can be managed.

Firstly, in modern MDE development environments, *models and metamodels are separated*. Metamodels are developed and distributed to users. Metamodels are installed, configured and combined to form a customised MDE development environment. Metamodel developers have no programmatic access to downstream instance models. Consequently, metamodel evolution occurs independently to model migration. First, the metamodel is evolved. Subsequently, the users of the metamodel find that their models are out-of-date and migrate their models.

Because of model and metamodel separation, existing techniques for managing co-evolution are either *developer-driven* (the metamodel developer devises an executable migration strategy, which is distributed to the metamodel user with the evolved metamodel) or *user-driven* (the metamodel user devises the migration strategy). In either case, it is always the metamodel user who performs migration on models that no longer conform to the metamodel.

Secondly, modern MDE development environments *implicitly enforce conformance*. A model is *bound* to its metamodel, typically by constructing a representation in the underlying programming language for each model element and data value. Frequently, binding is strongly-typed: each metamodel type is mapped to a corresponding type in the underlying programming language using

mappings defined by the metamodel. Consequently, MDE modelling frameworks do not permit changes to a model that would cause it to no longer conform to its metamodel. Loading a model that does not conform to its metamodel causes an error. In short, MDE modelling frameworks cannot be used to manage any model that does not conform to its metamodel.

Consequently, user-driven migration is always a manual process, in which models are migrated without using the modelling framework (which is unable to load non-conformant models). Executable migration strategies can only be used if they are specified with a tool that does not depend on the modelling framework to load the non-conformant models (and, at present, no such tool exists). Typically then, the metamodel user can only perform migration by editing the model directly, normally manipulating its underlying representation (e.g. XMI).

Furthermore, model migration must produce a model that conforms to the evolved metamodel because MDE modelling frameworks do not permit changes to a model that cause non-conformance. Therefore, model migration cannot be specified as a combination of co-evolution techniques with each performing some part of the migration.

Finally, a further consequence of implicitly enforced conformance is that models cannot be checked for conformance against any metamodel other than their own. Because conformance is always assumed, modern MDE development environments provide limited mechanisms for checking conformance, and typically provide no support for checking conformance to a metamodel other than the one used to construct the model. As will be discussed in Chapter ??, conformance checking can be used during metamodel installation to highlight potential conformance problems, and to determine whether migration is necessary for a model, permitting improved scalability for automatic model migration techniques.

### 4.2.2 Categorisation of Co-Evolution Techniques

Examples of co-evolution were analysed to discover and compare existing techniques for managing co-evolution. As discussed above, the separation of models and metamodels leads to two processes for co-evolution: *developer-driven* and *user-driven*. In this section, each process is described in more detail using a scenario. The challenges faced when using each process to perform co-evolution are discussed.

#### User-Driven Co-Evolution Techniques

In user-driven co-evolution, the metamodel user performs migration by locating non-conformant models and updating them to conform to the evolved metamodel. The metamodel developer might guide migration by providing a migration strategy to the metamodel user. Crucially, however, the migration strategy is not executable (e.g. it is written in prose). This is the key distinction between user-driven and developer-driven co-evolution. Only in the latter has the metamodel developer provided an executable model migration strategy.

In some cases, the metamodel user will not be provided with any migration strategy (executable or otherwise) from the metamodel developer. To perform migration, the metamodel user must determine which (if any) model elements no

longer conform to the evolved metamodel, and then decide how best to change non-conformant elements to re-establish conformance. This situation can be thought of as analogous to developing a legacy system<sup>3</sup>. Several examples of co-evolution required migration when no migration strategy had been provided by the metamodel developer.

**Scenario** The following scenario demonstrates user-driven co-evolution. Mark is developing an EMF metamodel. Members of his team, including Heather, will be using the metamodel to construct models. Heather installs the plug-in for Mark's metamodel and begins constructing models. Mark later identifies new requirements, changes the metamodel, builds a new version of the metamodel plug-in, and distributes it to Heather.

After several iterations of metamodel updates, Heather tries to load one of her older models, constructed using an earlier metamodel. When loading the older model, EMF reports an error indicating that the model no longer conforms to its metamodel. To load the older model, Heather must reinstall an old version of the metamodel plug-in that contains the metamodel to which the older model conforms. But even then, EMF will bind the older model to the old version of the metamodel, and not to the evolved metamodel. Worse still, Eclipse does not allow two versions of the same plug-in to co-exist in the same workspace, so the evolved metamodel cannot be used.

To avoid the loading errors in Heather's workspace, Mark could use an existing tool for managing co-evolution. However, all existing tools require extra effort from Mark. Operator-based approaches require him to evolve his metamodel using a specialised tool that records changes and applies co-evolutionary operators. Although metamodel matching approaches can automatically infer a migration strategy, they cannot be guaranteed to produce a migration strategy that capture the semantics that Mark intended.

Employing user-driven migration, Heather must trace and repair the loading error directly in XMI, the underlying format of the model. Human usability was not a key requirement for XMI [5] and, consequently, using XMI for migration is an unproductive and tedious task. Worse still, as EMF only reports the first problem encountered when deserialising an inconsistent model, re-establishing consistency in XMI is a slow, iterative process.

**Challenges** The above scenario highlights two of the major challenges faced when performing user-driven migration. Firstly, the underlying model representation is unlikely to be optimised for human usability. Together with limited support for conformance checking, migration by editing the underlying model representation is error prone and tedious. Secondly, installing a new version of a metamodel plug-in can affect the conformance of models and conformance problems are not reported to the user as part of the installation process. These challenges are elaborated on further in the final section of this chapter, which identifies research requirements.

## Developer-Driven Co-Evolution Techniques

further categorisation

---

<sup>3</sup>A system for which one has no documentation and whose authors have left the owning organisation.

### 4.3 Requirements Identification

When a metamodel change will affect conformance in only a small number of models, a metamodel developer may decide that the extra effort required to specify an executable migration strategy is too great, and prefer a user-driven co-evolution approach. Because no existing work has explored user-driven co-evolution and several of the co-evolution examples identified in Section X used a user-driven co-evolution approach, part of this thesis concentrated on exploring user-driven co-evolution.



# Bibliography

- [1] Eclipse. Model Development Tools project [online]. [Accessed 6 January 2008] Available at: <http://www.eclipse.org/modeling/mdt/>, 2009.
- [2] Martin Fowler. *Refactoring: improving the design of existing code*. Addison-Wesley, 1999.
- [3] Richard Gronback. Introduction to the Eclipse Graphical Modeling Framework. In *Proc. EclipseCon*, Santa Clara, California, 2006.
- [4] Anneke G. Kleppe, Jos Warmer, and Wim Bast. *MDA Explained: The Model Driven Architecture: Practice and Promise*. Addison-Wesley, 2003.
- [5] OMG. Human-Usable Textual Notation 1.0 Specification [online]. [Accessed 30 June 2008] Available at: <http://www.omg.org/technology/documents/formal/hutn.htm>, 2004.
- [6] OMG. Object Constraint Language 2.0 Specification [online]. [Accessed 30 June 2008] Available at: <http://www.omg.org/technology/documents/formal/ocl.htm>, 2006.
- [7] OMG. Unified Modelling Language 2.1.2 Specification [online]. [Accessed 30 June 2008] Available at: <http://www.omg.org/spec/UML/2.1.2/>, 2007.
- [8] OMG. Object Management Group home page [online]. [Accessed 30 June 2008] Available at: <http://www.omg.org>, 2008.
- [9] openArchitectureWare. openArchitectureWare Project Website [online]. [Accessed 30 June 2008] Available at: <http://www.eclipse.org/gmt/oaw/>, 2007.
- [10] Louis M. Rose, Dimitrios S. Kolovos, Richard F. Paige, and Fiona A.C. Polack. Model migration with flock. In *Proc. Fundamental Approaches to Software Engineering*, 2010.
- [11] Louis M. Rose, Richard F. Paige, Dimitrios S. Kolovos, and Fiona A.C. Polack. An analysis of approaches to model migration. In *Proc. Joint MoDSE-MCCM Workshop*, 2009.
- [12] Dag I.K. Sjøberg. Quantifying schema evolution. *Information & Software Technology*, 35(1):35–44, 1993.
- [13] Andrew Watson. A brief history of MDA. *Upgrade*, 9(2), 2008.