# Evolution in Model-Driven Engineering

Louis M. Rose

August 2, 2010

# Contents

# Chapter 3

# Literature Review

This chapter provides a review and critical analysis of existing work on software evolution and identifies potential research directions. The principles of software evolution are discussed in Section 3.1, while Section 3.2 reviews the ways in which evolution is identified, analysed and managed in a range of fields, including relational databases, programming languages, and model-driven development environments. From the reviewed literature, Section 3.3 synthesises research challenges for software evolution in the context of MDE, highlighting those to which this thesis contributes, and elaborates on the research method used in this thesis.

## 3.1   Software Evolution Theory

Software evolution is an important facet of software engineering. Studies [Erlikh 2000, Moad 1990] suggest that the evolution of software can account for as much as 90% of a development budget. Such figures are sometimes described as uncertain [Sommerville 2006, ch. 21], primarily because the term evolution is not used consistently. Nonetheless, there is a corpus of software evolution research, and publications in this area have existed since the 1960s (e.g. [Lehman 1969]).

The remainder of this section introduces software evolution terminology and discusses three research areas that relate to software evolution: refactoring, design patterns and traceability. Refactoring concentrates on improving the structure of existing systems, design patterns on best practices for software design, and traceability for recording and analysing the lifecycle of software artefacts. Each area provides a common vocabulary for discussing software design and evolution. There is an abundance of research in these areas, including seminal works on refactoring by [Opdyke 1992] and [Fowler 1999]; and on design patterns by [Alexander *et al.* 1977] and [Gamma *et al.* 1995].

### 3.1.1   Categories of Software Evolution

[Sjøberg 1993] identifies reasons for software evolution, which include addressing changing requirements, adapting to new technologies, and architectural restructuring. These reasons are the motivations for three common types of software evolution [Sommerville 2006, ch. 21]:

- **Corrective evolution** takes place when a system exhibiting unintended or faulty behaviour is corrected. Alternatively, corrective evolution may be used to adapt a system to new or changing requirements.

- **Adaptive evolution** is employed to make a system compatible with a change to platforms or technologies that underpin its implementation.

- **Perfective evolution** refers to the process of improving the internal quality of a system, while preserving the behaviour of the system.

The remainder of this section adopts this categorisation for discussing software evolution literature. Refactoring (discussed in Section 3.1.2), for instance, is one way in which perfective evolution can be realised.

Many activities are used for identifying and managing software evolution. [Winkler & Pilgrim 2009] highlight the importance of *impact analysis* (for reasoning about the effects of evolution) and *change propagation* (for updating one artefact in response to a change made to another). In addition, [Sommerville 2006] notes that *reverse engineering* (analysing existing development artefacts to extract information) and *source code translation* (rewriting code to use a more suitable technology, such as a different programming language) are also important software evolution activities. MDE facilitates portable software, for example by prescribing platform-independent and platform-specific models (as discussed in Section 2.1.3), and as such source code translation is arguably less relevant to MDE than to traditional software engineering. Because MDE seeks to capture the essence of the software in models, reverse engineering information from, for example, code is also less likely to be relevant to MDE than to tradition software engineering. Consequently, this thesis focuses on impact analysis and change propagation.

### 3.1.2   Refactoring

[Mens & Tourwé 2004] report "an urgent need for techniques that reduce software complexity by incrementally improving the internal software quality." Refactoring was first described by [Opdyke 1992] and is "the process of changing a software system in such a way that it does not alter the external behaviour of the code yet improves its internal structure" [Fowler 1999, pg. xvi]. Refactoring plays a significant role in the evolution of software systems – a recent study of five open-source projects showed that over 80% of changes were refactorings [Dig & Johnson 2006b].

Typically, refactoring literature concentrates on three primary activities in the refactoring process: *identification* (where should refactoring be applied, and which refactorings should be used?), *verification* (has refactoring preserved behaviour?) and *assessment* (how has refactoring affected other qualities of the system, such as cohesion and efficiency?).

In the foreword to [Fowler 1999], Beck describes an informal means for identifying the need for refactoring, termed *bad smells*: "structures in the code that suggest (sometimes scream for) the possibility of refactoring.". Tools and semi-automated approaches have also been devised for refactoring identification, such as Daikon [Kataoka *et al.* 2001], which detects program invariants that may indicate the possibility for refactoring. Clone analysis tools have been employed for identifying refactorings that eliminate duplication [Balazinska *et al.* 2000,

Ducasse *et al.* 1999]. The types of refactoring being performed may vary over different domains. For example, Buck[1] describes refactorings, such as "Skinny Controller, Fat Model", particular to the Ruby on Rails web framework [37-Signals 2008].

MOF [OMG 2008a], discussed in Section 2.1.3, provides a standard notation for describing the abstract syntax of metamodels. As MOF re-uses many concepts from UML class diagrams (which are used to describe the structure of object-oriented systems), object-oriented refactorings can be applied to metamodels defined using MOF. However, no standard means has yet been defined for attaching semantics to modelling language constructs. When a metamodel is defined without a rigorous semantics, refactoring of the sort applied to OO code does not seem to be directly applicable. (In particular, drawing parallels to existing approaches for the verification and assessment activities of refactoring seems difficult). Regardless, refactoring catalogues, such as [Fowler 1999], might influence the way in which model evolution is recorded, due to the clarity and conciseness of their format. This is discussed further in Section 3.1.3.

Since 2006, Dig has been studying the refactoring of systems that are developed by combining components, possibly developed by different organisations. [Dig & Johnson 2006b] reports a survey used to identify and categorise the changes made to five components that are known to have been re-used often, with the hypothesis that a significant number of the changes could be classified as behaviour-preserving (i.e. refactorings). By using examples from the survey, [Dig *et al.* 2006] devises an algorithm for automatically detecting refactorings to a high degree of accuracy (over 85%). The algorithm was then utilised in tools for (1) replaying refactorings to perform migration of client code following breaking changes to a component [Dig & Johnson 2006a], and (2) versioning object-oriented programs using a refactoring-aware configuration management system [Dig *et al.* 2007]. The latter facilitated better understanding of program evolution, and the refinement of the refactoring detection algorithm.

### 3.1.3 Patterns and anti-patterns

A *design pattern* identifies a commonly occurring design problem and describes a re-usable solution to that problem. Related design patterns are combined to form a *pattern catalogue* – such as for object-oriented programming [Gamma *et al.* 1995] or enterprise applications [Fowler 2002]. A pattern description comprises at least a name, overview of the problem, and details of a common solution [Brown *et al.* 1998]. Depending on the domain, further information may be included in the pattern description (such as a classification, a description of the pattern's applicability and an example usage).

Design patterns can be thought of as describing objectives for improving the internal quality of a system (perfective software evolution). [Kerievsky 2004] provides a practical guide that describes how software can be refactored towards design patterns to improve its quality. Studying the way in which experts perform perfective software evolution can lead to devising best practices, sometimes in the form of a pattern catalogue, such as the object-oriented refactorings described in [Fowler 1999].

[Alexander *et al.* 1977] first used design patterns when devising a pattern catalogue for town planning. [Beck & Cunningham 1989] later adapted the

---

[1]In a keynote address to the First International Ruby on Rails Conference (RailsConf), May 2007, Portland, Oregon, United States of America.

work of Alexander for software architecture, by specifying a pattern catalogue for designing user-interfaces. Utilising pattern catalogues allowed the software industry to "reuse the expertise of experienced developers to repeatedly train the less experienced." [Brown *et al.* 1998, pg. 10]. [Rising 2001, pg. xii] summarises the usefulness of design patterns: "Patterns help to define a vocabulary for talking about software development and integration challenges; and provide a process for the orderly resolution of these challenges."

Anti-patterns are an alternative literary form for describing patterns of a software architecture [Brown *et al.* 1998]. Rather than describe patterns that have often been observed in successful architectures, they describe those which are present in unsuccessful architectures. Essentially, an anti-pattern is a pattern in an inappropriate context, which describes a problematic solution to a frequently encountered problem. The (anti-)pattern catalogue may include alternative solutions that are known to yield better results (termed "refactored solutions" by [Brown *et al.* 1998]). Catalogues might also consider the reasons why (inexperienced) developers might select an anti-pattern. Brown notes that "patterns and anti-patterns are complementary" [Brown *et al.* 1998, pg. 13]; both are useful in providing a common vocabulary for discussion of system architectures and in educating less experienced developers.

### 3.1.4   Traceability

A software development artefact rarely evolves in isolation. Changes to one artefact cause and are caused by changes to other artefacts (e.g. object code is recompiled when source code changes, source code and documentation are updated when requirements change). Hence, traceability – the ability to describe and follow the life of software artefacts [Winkler & Pilgrim 2009, Lago *et al.* 2009] – is closely related to and facilitates software evolution.

Historically, traceability is a branch of requirements engineering, but increasingly traceability is used for artefacts other than requirements [Winkler & Pilgrim 2009]. Because MDE prescribes automated transformation between models, traceability is also researched in the context of MDE. The remainder of this section discusses traceability principles focussing on the relationship between traceability and software evolution, while Section 3.2.4 reviews the traceability literature that relates to MDE.

Traceability is facilitated by *traceability links*, which document the dependencies, causalities and influences between artefacts. Traceability links are established by hand or by automated analysis of artefacts. In MDE environments, some traceability links can be automatically inferred because the relationships between some types of artefact are specified in a structured manner (for example, as a model-to-model transformation).

Traceability links are defined between artefacts at the same level of abstraction (horizontal links) and at different levels of abstraction (vertical links). Uni-directional traceability links are navigated either *forwards* (away from the dependent artefact) or *backwards* (toward the dependent artefact). Figure 3.1 summaries these categories of traceability link.

The traceability literature uses inconsistent terminology. This thesis adopts the same terminology as [Winkler & Pilgrim 2009]: *traceability* is the ability to describe and follow the life of software artefacts; *traceability links* are the relationships between software artefacts.
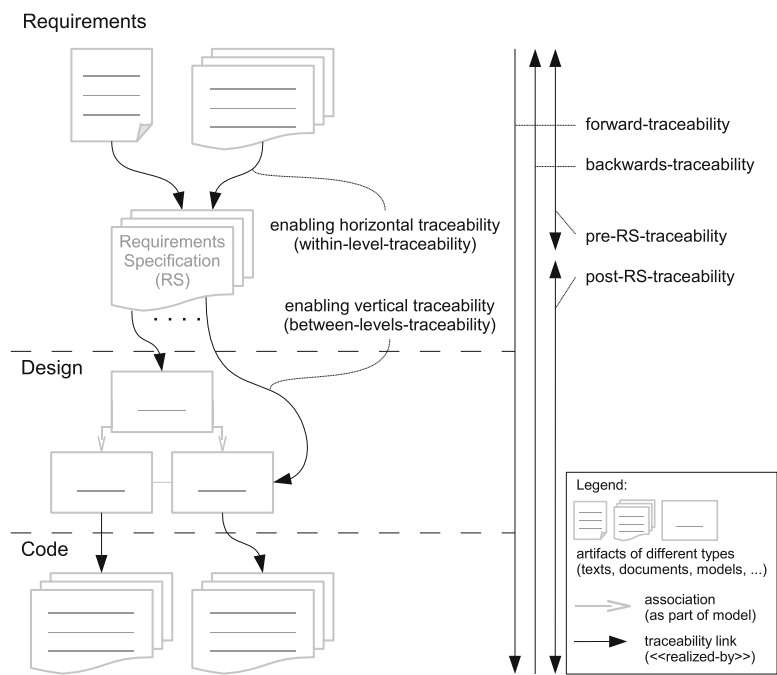
Figure 3.1: Categories of traceability link [Winkler & Pilgrim 2009].

Traceability supports software evolution activities, such as impact analysis (discovering and reasoning about the effects of a change) and change propagation (updating impacted artefacts following a change to an artefact). Moreover, automated software evolution is facilitated by programmatic access to traceability links.

Current approaches for traceability-supported software evolution use *triggers* and *events*. Each approach proposes mechanisms for detecting triggers (changes to artefacts) and for notifying dependent artefacts of events (the details of a change). Existing approaches vary in the extent to which they can automatically update dependent artefacts. The approaches described in [Chen & Chou 1999, Cleland-Huang *et al.* 2003] report inconsistencies and do not perform automatic updates, while [Aizenbud-Reshef *et al.* 2005, Costa & Silva 2007] propose reactive approaches for guided or fully automatic updates. Section 3.2.4 provides a more thorough discussion and critical analysis of event-based approaches for impact analysis and change propagation in the context of MDE.

To remain accurate and hence useful, traceability links must be updated as a system evolves. Although most existing approaches to traceability are "not well suited to the evolution of [traceability] artefacts" [Winkler & Pilgrim 2009, pg. 24], there is some work in this area. For example, [Mäder *et al.* 2008] describe a development environment that records changes to artefacts, comparing the changes to a catalogue of built-in patterns. Each pattern provides an executable specification for updating traceability links.

Software evolution and traceability are entangled concerns. Traceability facilitates software evolution activities such as impact analysis and change propagation. Traceability is made possible with consistent and accurate traceability links. Software evolution can affect the relationships between artefacts (i.e. the traceability links) and hence software evolution techniques are applied to ensure that traceability links remain consistent and accurate.

## 3.2   Software Evolution in Practice

Using the principles of software evolution described above, this section examines the ways in which evolution is identified, managed and analysed in a variety of settings, including programming languages grammarware, relational database management system and MDE.

### 3.2.1   Programming Language Evolution

Programming language designers often attempt to ensure that legacy programs continue to conform to new language specifications. For example, [Cervelle *et al.* 2006] highlights that the Java [Gosling *et al.* 2005] language designers are reluctant to introduce new keywords (as identifiers in legacy programs could then be mistakenly recognised as instances of the new keyword).

Although designers are cautious about changing programming languages, evolution does occur. In this section, two examples of the ways in which programming languages have evolved are discussed. The vocabulary used to describe the scenarios is applicable to evolution of MDE artefacts. Furthermore, MDE sometimes involves the use of general-purpose modelling languages, such

as UML [OMG 2007a]. The evolution of general-purpose modelling languages may be similar to that of general-purpose programming languages.

### Reduction

Mapping language abstractions to executable concepts can be complicated. Therefore, languages are sometimes evolved to simplify the implementation of translators (compilers, interpreters, etc). It seems that this type of evolution is more likely to occur when language design is a linear process (with a reference implementation occurring after design), and in larger languages.

[Backus 1978] identifies some simplification during FORTRAN's evolution: originally, FORTRAN's `DO` statements were awkward to compile. The semantics of `DO` were simplified such that more efficient object code could be generated from them. Essentially, the simplified `DO` statement allowed linear changes to index statements to be detected (and optimised) by compilers.

The removal of the `RELABEL` construct (which facilitated more straightforward indexing into multi-dimensional arrays) from the FORTRAN language specification [Backus 1978] is a further example of reduction.

### Revolution

Developers often form best practices for using languages. Design patterns are one way in which best practices may be communicated with other developers. Incorporating existing design patterns as language constructs is one approach to specifying a new language (e.g. [Bosch 1998]).

Lisp makes idiomatic some of the Fortran List Processing Language (FLPL) design patterns. For example, [McCarthy 1978] describes the awkwardness of using FLPL's `IF` construct, and the way in which experienced developers would often prefer to define a function of the form `XIF(P, T, F)` where `T` was executed iff `P` was true, and `F` was executed otherwise. However, such functions had to be used sparingly, as all three arguments would be evaluated due to the way in which FORTRAN executed function calls. McCarthy [McCarthy 1978] defined a more efficient semantics, wherein `T` (`F`) was only evaluated when `P` was true (false). Because FORTRAN programs could not express these semantics, McCarthy's new construct informed the design of Lisp. Lazy evaluation in functional languages can be seen as a further step on this evolutionary path.

## 3.2.2 Schema Evolution

This section reviews schema evolution research. Work covering the evolution of XML and database schemata is considered. Both types of schema are used to describe a set of concepts (termed the *universe of discourse* in database literature). Schema designers decide which details of their domain concepts to describe; their schemata provide an abstraction containing only those concepts which are relevant [Elmasri & Navathe 2006, pg. 30]. As such, schemata in these domains may be thought of as analogous to metamodels – they provide a means for describing an abstraction over a phenomenon of interest. Therefore, approaches to identifying, analysing and performing schema evolution are directly relevant to the evolution of metamodels in MDE. However, the patterns of

evolution commonly seen in database systems and with XML may be different
to those of metamodels because evolution can be:

- **Domain-specific**: Patterns of evolution may be applicable only within a
  particular domain (e.g. normalisation in a relational database).

- **Language-specific**: The way in which evolution occurs may be influ-
  enced by the language (or tool) used to express the change. (For example,
  some implementations of SQL may not have a `rename relation` com-
  mand, so alternative means for renaming a relation must be used).

Many of the published works on schema evolution share a similar method,
with the aim of defining a taxonomy of evolutionary operators. Schema main-
tainers are expected to employ these operators to change their schemata. This
approach is used heavily in the XML schema evolution community, and was the
sole strategy encountered [Guerrini *et al.* 2005, Kramer 2001, Su *et al.* 2001].
Similar taxonomies have been defined for schema evolution in relational database
systems (e.g. in [Banerjee *et al.* 1987, Edelweiss & Freitas Moreira 2005]), but
other approaches to evolution are also prevalent. One alternative, proposed in
[Lerner 2000], is discussed in depth, along with a summary of other work.

**XML Schema Evolution**

XML provides a specification for defining mark-up languages. XML documents
can reference a schema, which provides a description of the ways in which the
concepts in the mark-up should relate (i.e. the schema describes the syntax of
the XML document). Prior to the definition of the XML Schema specification
[W3C 2007a] by the W3C [W3C 2007b], authors of XML documents could use a
specific Document Type Definition (DTD) to describe the syntax of their mark-
up language. XML Schemata provide a number of advantages over the DTD
specification:

- XML Schemata are defined in XML and may, therefore, be validated
  against another XML Schema. DTDs are specified in another language
  entirely, which requires a different parser and different validation tools.

- DTDs provide a means for specifying constraints only on the mark-up
  language, whereas XML Schemata may also specify constraints on the
  data in an XML document.

Work on the evolution of the structure of XML documents is now dis-
cussed. [Guerrini *et al.* 2005] concentrate on changes made to XML Schema,
while [Kramer 2001] focuses on DTDs.

[Guerrini *et al.* 2005] propose a set of primitive operators for changing XML
schemata. They show this set to be both sound (application of an operator
always results in a valid schema) and complete (any valid schema can be pro-
duced by composing operators). Their classification also details those operators
that are 'validity-preserving' (i.e. application of the operator produces a schema
that does not require its instances to be migrated). Guerrini et al. show that
the arguments of an operator can influence whether it is validity-preserving.
For example, inserting an element is validity-preserving when inclusion of the
element is optional for instances of the schema. In addition to soundness and

completeness, minimality is another desirable property in a taxonomy of primitive operators for performing schema evolution [Su *et al.* 2001]. To complement a minimal set of primitives, and to improve the conciseness with which schema evolutions can be specified, [Guerrini *et al.* 2005] propose a number of 'high-level' operators, which comprise two or more primitive operators.

[Kramer 2001] provides another taxonomy of primitives for XML schema evolution. To describe her evolution operators, Kramer uses a template, which comprises a name, syntax, semantics, preconditions, resulting DTD changes and resulting data changes section for each operator. This style is similar to a pattern catalogue, but Kramer does not provide a context for her operators (i.e. there are no examples that describe when the application of an operator may be useful). Kramer utilises her taxonomy in a repository system, Exemplar, for managing the evolution of XML documents and their schemata. The repository provides an environment in which the variation of XML documents can be managed. However, to be of practical use, Exemplar would benefit from integration with a source code management system (to provide features such as branching, and version merging).

As noted in [Pizka & Jürgens 2007], the approaches described in [Kramer 2001, Su *et al.* 2001, Guerrini *et al.* 2005] are complete in the sense that any valid schema can be produced, but do not allow for arbitrary updates of the XML documents in response to schema changes. Hence, none of the approaches discussed in this section ensure that information contained in XML documents is not lost.

**Relational Database Schema Evolution**

Defining a taxonomy of operators for performing schema updates is also common for supporting relational database schema evolution (e.g. [Edelweiss & Freitas Moreira 2005, Banerjee *et al.* 1987]). However, [Lerner 2000] highlights problems that arise when performing data migration after these taxonomies have been used to specify schema evolution:

> "There are two major issues involved in schema evolution. The first issue is understanding how a schema has changed. The second issue involves deciding when and how to modify the database to address such concerns as efficiency, availability, and impact on existing code. Most research efforts have been aimed at this second issue and assume a small set of schema changes that are easy to support, such as adding and removing record fields, while requiring the maintainer to provide translation routines for more complicated changes. As a result, progress has been made in developing the backend mechanisms to convert, screen, or version the existing data, but little progress has been made on supporting a rich collection of changes" [Lerner 2000, pg. 84].

Fundamentally, [Lerner 2000] believes that any taxonomy of operators for schema evolution is too fine-grained to capture the semantics intended by the schema developer, and therefore cannot be used to provide automated migration: [Lerner 2000] states that existing taxonomies are concerned with the "editing process rather than the editing result". Furthermore, Lerner believes that developing such a taxonomy creates a proliferation of operators, increasing the

complexity of specifying migration. To demonstrate, Lerner considers moving a field from one type to another in a schema. This could be expressed using two primitive operators, `delete_field` and `add_field`. However, the semantics of a `delete_field` command likely dictate that the data associated with the field will be lost, making it unsuitable for use when specifying that a type has been moved. The designer of the taxonomy could introduce a `move_field` command to solve this problem, but now the maintainer of the schema needs to understand the difference between the two ways in which moving a type can be specified, and carefully select the correct one. Lerner provides other examples which elucidate this issue (such as introducing a new type by splitting an existing type). Even though [Lerner 2000] highlights that a fine-grained approach may not be the most suitable for specifying schema evolution, other potential uses for a taxonomy of evolutionary operators (such as being used as a common vocabulary for discussing the restructuring of a schema) are not discussed.

[Lerner 2000] proposes an alternative to operator-based schema evolution in which two versions of a schema are compared to infer the schema changes. Using the inferred changes, migration strategies for the affected data can be proposed. [Lerner 2000] presents algorithms for inferring changes from schemata and performing both automated and guided migration of affected data. By inferring changes, developers maintaining the schema are afforded more flexibility. In particular, they need not use a domain-specific language or editor to change a schema, and can concentrate on the desired result, rather than how best to express the changes to the schema in the small. Furthermore, algorithms for inferring changes have use other than for migration (e.g. for semantically-aware comparison of schemata, similar to that provided by a refactoring-aware *source code management system*, such as [Dig *et al.* 2007]). Comparison of two schema versions might suggest more than one feasible strategy for updating data, and [Lerner 2000] does not propose a mechanism for distinguishing between feasible alternatives.

[Vries & Roddick 2004] propose the introduction of an extra layer to the architecture typical of a relational database management system. They demonstrate the way in which the extra layer can be used to perform migration subsequent to a change of an attribute type. The layer contains (mathematical) relations, termed *mesodata*, that describe the way in which an old value (data prior to migration) maps to one or more new values (data subsequent to migration). These mappings are added to the mesodata by the developer performing schema updates, and are used to semi-automate migration. It is not clear how this approach can be applied when schema evolution is not an attribute type change.

In the O2 database [Ferrandina *et al.* 1995], schema updates are performed using a small domain-specific language. Modification constructs are used to describe the changes to be made to the schema. To perform data migration, O2 provides conversion functions as part of its modification constructs. Conversion functions are either user-defined or default (pre-defined). The pre-defined functions concentrate on providing mappings for attributes whose types are changed (e.g. from a double to an integer; from a set to a list). Additionally, conversion functions may be executed in conjunction with the schema update, or they may be deferred, and executed only when the data is accessed through the updated schema. Ferrandina et al. observe that deferred updates may prevent unnecessary downtime of the database system. Although the approach outlined in

[Ferrandina *et al.* 1995] addresses the concern that "approaches to coping with schema evolution should be concerned with the editing result rather than the editing process" [Lerner 2000], there is no support for some types of evolution such as moving an attribute from one relation to another.

### 3.2.3 Grammar Evolution

[Klint *et al.* 2003] call for an engineering approach to producing grammarware (grammars and software that depends on grammars, such as parsers and program convertors). The grammarware engineering approach envisaged by Klint et al. is based on best practices and techniques, which they anticipate will be derived from addressing open research challenges. Klint et al. identify seven key questions for grammarware engineering, one of which relates to grammar evolution: "How does one systematically transform grammatical structure when faced with evolution?" [Klint *et al.* 2003, pg. 334].

Between 2001 and 2005, Ralf Lämmel (an author of [Klint *et al.* 2003]) and his colleagues at Vrije Universiteit published several important papers on grammar evolution. [Lämmel 2001] proposes a taxonomy of operators for semi-automatic grammar refactoring and demonstrates their usefulness in recovering the formal specifications of undocumented grammars (such as VS COBOL II in [Lämmel & Verhoef 2001]) and in specifying generic refactorings [Lämmel 2002].

The work of Lämmel et al. focuses on grammar evolution for refactoring or for *grammar recovery* (corrective evolution in which a deviation from a language reference is removed), but does not address the impact of grammar evolution on corresponding programs or grammarware. For instance, when a grammar changes, updates are potentially required both to programs written in that grammar and to tools that parse, generate or otherwise manipulate programs written in that grammar.

[Pizka & Jürgens 2007] recognise and seek to address the challenge of grammar-program co-evolution. Pizka and Juergens believe that most grammars evolve over time and that, without tool support, co-evolution is a complex, time-consuming and error prone task. To this end, [Pizka & Jürgens 2007] proposes Lever, a language evolution tool, which defines and uses operators for changing grammars (and programs) in an approach that is inspired by [Lämmel 2001].

Compared to the taxonomy in [Lämmel 2001], Lever can be used to manage the evolution of grammars, programs and the co-evolution of grammars and programs, and the taxonomy defined by Lämmel et al. can be used only to manage grammar evolution. However, as a consequence, Lever sacrifices the formal preservation properties of the taxonomy defined by Lämmel et al.

### 3.2.4 Evolution of MDE Artefacts

As discussed in Chapter 1, the evolution of development artefacts during MDE inhibits the productivity and maintainability of model-driven approaches for constructing software systems. Mitigating the effects of evolution on MDE is an open research topic, to which this thesis contributes.

This section discusses literature that explores the evolution of development artefacts used when performing MDE. [Deursen *et al.* 2007] highlight that evolution in MDE is complicated, because it spans multiple dimensions. In particular, there are three types of development artefact specific to MDE: models,

metamodels, and specifications of model management tasks[2]. A change to one type of artefact can affect other artefacts (possibly of a different type).

[Sprinkle & Karsai 2004] highlights that the evolution of an artefact can appear to be either *syntactic* or *semantic*. In the former, no information is known about the intention of the the evolutionary change. In the latter, a lack of detailed information about the semantics of evolution can reduce the extent to which change propagation can be automated. For example, consider the case where a class is deleted from a metamodel. The following questions typically need to be answered to facilitate evolution:

- Should subtypes of the deleted class also be removed? If not, should their inheritance hierarchy be changed? What is the correct type for references that used to have the type of the deleted class?

- Suppose that the evolving metamodel was the target of a previous model-to-model transformation. Should the data that was previously transformed to instances of the deleted class now be transformed to instances of another metamodel class?

- What should happen to instances of the deleted metamodel class? Perhaps they should be removed too, or perhaps their data should be migrated to new instances of another class.

Tools that recognise only syntactic evolution tend to lack the information required for full automation of evolution activities. Furthermore, tools that focus only upon syntax cannot be applied in the face of additive changes [Gruschko *et al.* 2007]. There are complexities involved in recording the semantics of software evolution. For example, the semantics of an impacted artefact need not always be preserved: this is often the case in corrective evolution.

Notwithstanding the challenges described above, MDE has great potential for managing software evolution and automating software evolution activities, particularly because of model transformations (Section **??**). Approaches for managing evolution in other fields, described above, must consider the way in which artefacts are updated when changes are propagated from one artefact to another. Model transformation languages already fulfil this role in MDE. In addition, model transformations provide a (limited) form of traceability between MDE artefacts, which can be used in impact analysis.

This section focuses on the three types of evolution most commonly discussed in model-driven engineering literature. *Model refactoring* is used to improve the quality of a model without changing its functional behaviour. *Model synchronisation* involves updating a model in response to a change made in another model, usually by executing a model-to-model transformation. *Model-metamodel co-evolution* involves updating a model in response to a change made to a metamodel. This section concludes by reviewing existing techniques for visualising model-to-model transformation and assessing their usefulness for understanding evolution in the context of MDE.

---

[2]Some examples of model management tasks include model-to-model transformation, model-to-text transformation, model validation, model merging and model comparison.

**Model Refactoring**

Refactoring (Section 3.1.2) is a perfective software evolution activity in which the structure and representation of a system is improved without changing its functional behaviour. Refactoring has been studied in the context of MDE because refactoring can be domain-specific (e.g. normalisation in relational databases). Model refactoring languages allow metamodel developers to capture commonly occurring refactoring patterns and provide their users with model editors that support automatic refactoring.

In model transformation terminology (discussed in Section **??**), a refactoring is an *endogenous*, *in-place* transformation. Refactorings are applied to an artefact (e.g. model, code) producing a semantically equivalent artefact, and hence an artefact that conforms to the same rules and structures as the original. Because refactorings are used to improve the structure of an existing artefact, the refactored artefact typically replaces the original. Endogenous, in-place transformation languages, suitable for refactoring, are described in [Biermann *et al.* 2006, Porres 2003] (which propose declarative approaches based on graph theory) and in [Kolovos *et al.* 2007] (which proposes mixing declarative and imperative constructs).

There are similarities between the structures defined in the MOF metamodelling language and in object-oriented programming languages. For the latter, refactoring pattern catalogues exist (such as [Fowler 1999]), which might usefully be applied to modelling languages. [Moha *et al.* 2009] provides a notation for specifying refactorings for MOF and UML models and Java programs in a generic (metamodel-independent) manner. Because MOF, UML and the Java language share some concepts with the same semantics (such as classes and attributes), [Moha *et al.* 2009] show that refactorings can be shared among them, but only consider 3 of the object-oriented refactorings identified in [Fowler 1999]. To more thoroughly understand metamodel-independent refactoring, a larger number of refactorings and languages should be explored.

Abstraction is a fundamental benefit of MDE (Section **??**). Defining a domain-specific language is one way in which abstraction can be realised for MDE (Section **??**). In addition to tools for defining modelling languages, generating model editors and performing model transformation, model-driven development environments might benefit from mechanisms for defining domain-specific refactorings. In particular, metamodel developers may wish to document common patterns of evolution, perhaps in an executable format.

Eclipse, an extensible development environment, provides a library for building development tools for textual languages, LTK (language toolkit) [**?**]. LTK allows developers to specify – in Java – refactorings for their language, which can be invoked via the language editor. LTK makes no assumptions on the way in which languages will be structured, and as such refactoring code that operates on models must interact with the modelling framework directly.

The Epsilon Wizard Language (EWL) [Kolovos *et al.* 2007] is a model transformation language tailored for the specification of model refactorings. EWL is built atop Epsilon and its object language (EOL), which can query, update and navigate models represented in a diverse range of modelling technologies (Section 2.3.2). Consequently, EWL, unlike LTK, abstracts over modelling frameworks.

[Arendt *et al.* 2009] present EMF Refactor, comparing it with EWL and the

LTK by specifying a refactoring on a UML model. EMF Refactor, like EWL, contributes a model transformation language tailored for refactoring. In contrast to EWL, EMF Refactor has a visual (rather than textual) syntax, and is based on graph transformation concepts. Figure 3.2 shows the "Change attribute to association end" refactoring for the UML metamodel in EMF Refactor. The left-hand side of the refactoring rule (Figure 3.2(a)) matches a Class whose owned attributes contains a Property whose type has the same name as a Class. The right-hand side of the rule (Figure 3.2(b)) introduces a new Association, whose member end is the Property matched in the left-hand side of the rule. Due to the visual syntax, EMF Refactor might be usable only with modelling technologies based on MOF (which has a graphical concrete-syntax based on UML class diagrams). From [Arendt *et al.* 2009], it is not clear to what extent EMF Refactor can be used with modelling technologies other than EMF.



(a) Left-hand side matching rule.
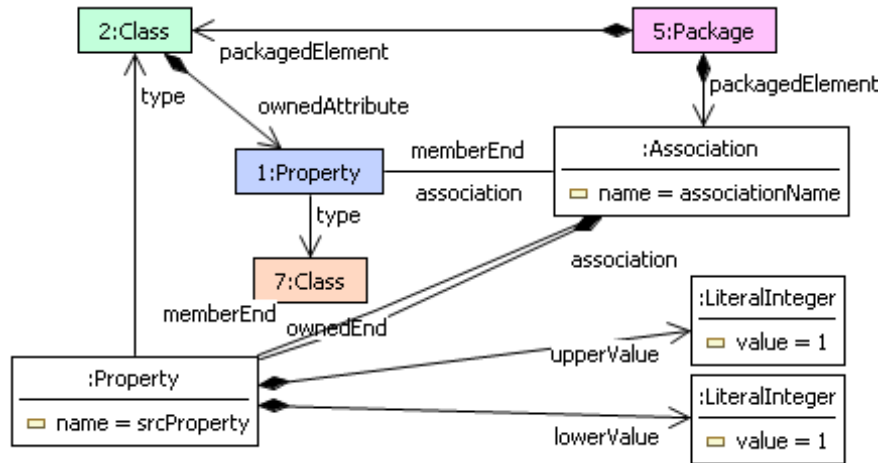


(b) Right-hand side production rule.

Figure 3.2: Attribute to association end refactoring in EMF Refactor. Taken from [Arendt *et al.* 2009].

[Kolovos *et al.* 2007] and [Arendt *et al.* 2009] focus on refactoring a model in isolation. Neither approach can be used to specify *inter-model refactorings*, which impact more than one model at once. The Eclipse Java Development Tools support refactorings of Java code that update many source-code artefacts

at once: for example, renaming a class in one source file updates references to that class in other source files. In the context of MDE, support for inter-model refactoring would facilitate a greater degree of model modularisation, regarded by [Kolovos *et al.* 2008c] as a solution to scalability, one of the challenges faced by MDE.

According to [Mens *et al.* 2007], "research in model refactoring is still in its infancy." Mens et al. identify formalisms for investigating the feasibility and scalability of model refactoring. In particular, Mens et al. suggest that meaning-preservation (an objective of refactoring, as discussed in Section 3.1.2) can be checked by evaluating OCL constraints, behavioural models or downstream program code.

### Model Synchronisation

Changes made to development artefacts may require the *synchronisation* of related artefacts (models, code, documentation). Traceability links (which capture the relationships of software artefacts) facilitate synchronisation. This section discusses the way in which change propagation is approached in the literature, which typically involves using an incremental style of transformation. Work that addresses more fundamental aspects of model synchronisation, such as capturing trace links and performing impact analysis are also discussed. Finally, synchronisation between models and text and between models and trace links is also considered.

**Incremental Transformation** Many model synchronisation approaches extend or instrument existing model-to-model transformation languages. Declarative transformation languages lend themselves to the specification of bi-directional transformations (which [Fritzsche *et al.* 2008] describe as traceability-by-design) and *incremental transformations*, a style of model transformation that facilitates incremental updates of the target model. In fact, most model synchronisation literature focuses on incremental transformation.

Incremental transformation is most often achieved in one of two ways. Because model-to-model transformation is used to generate one or more target models from one or more source models, when a source model changes, the model-to-model transformation can be invoked to completely re-generate the target models. [Hearnden *et al.* 2006] call this activity *re-transformation*, and propose an alternative approach, *live transformation*, in which the transformation context is persistent. Figure 3.3 illustrates the differences between re transformation and live transformation, showing the evolution of source and target models on the left-hand and right-hand sides, respectively, and the transformation context in the middle. Live transformation facilitates change propagation from the source to the target models without completely re-generating the target models and is therefore a more efficient approach. As well as in [Hearnden *et al.* 2006], live transformation is used to achieve incremental transformation in [Ráth *et al.* 2008] and [Tratt 2008].

Primarily, incremental transformation has been used to address the scalability of model transformations. For large models, transformation execution time has been shown to be significantly reduced by using incremental transformation [Hearnden *et al.* 2006]. However, [Kolovos *et al.* 2008c] suggests that scalability should be addressed not only by attempting to develop techniques for increasing

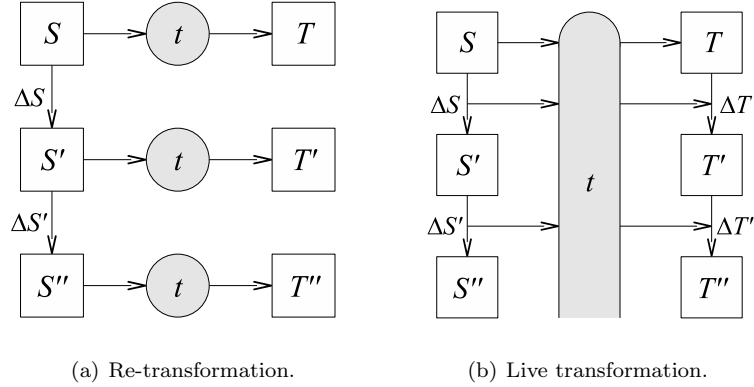(a) Re-transformation.                    (b) Live transformation.

Figure 3.3:    Approaches  to  incremental  transformation.    Taken  from
[Hearnden *et al.* 2006].

the speed of model transformation, but also by providing principles, practices
and tools for building models that are less monolithic and more modular. For
this end, model synchronisation research that focuses solely on increasing scal-
ability is unhelpful and should instead focus on improving maintainability in
conjunction with – or rather than – scalability.

Model synchronisation and incremental transformation can be applied to de-
couple models and facilitate greater modularisation, although this is not com-
monly discussed in the literature. [Fritzsche *et al.* 2008] describe an automated,
model-driven approach to performance engineering. Fritzsche et al. contribute a
transformation that produces, from any UML model, a model for which perfor-
mance characteristics can be readily analysed. The relationships between UML
and performance model artefacts are recorded using traceability links. The re-
sults of the performance analysis are later fed back to the UML model using
an incremental transformation made possible by the traceability links. Using
this approach, performance engineers can focus primarily on the performance
models, while other engineers are shielded from the low-level detail of the per-
formance analysis. As such, Fritzsche et al. show that two different modelling
concerns can be separated and decoupled, yet remain cohesive via the applica-
tion of model synchronisation.

**Towards automated model synchronisation**    Some existing work provides
a foundation for automating model synchronisation activities. Theoretical as-
pects of the traceability literature were reviewed in Section 3.1.4, and explored
the automated activities that traceability facilitates, such as impact analysis
and change propagation. This section now analyses the traceability research
in the context of model-driven engineering and focuses on the way in which
traceability facilitates the automation of model synchronisation activities.

Aside from live transformation, other techniques for capturing trace links be-
tween models have been reported. Enriching a model-to-model transformation
with traceability information is discussed in [Jouault 2005], which contributes a
generic higher-order transformation for this purpose. Given a transformation,
the generic higher-order transformation adds transformation rules that produce

a traceability model. In contrast to the genericity of the approach described in [Jouault 2005], [Drivalos *et al.* 2008] propose domain-specific traceability meta-models for richer traceability link semantics. Further research is required to assess the requirements of automated model synchronisation tools and to select appropriate traceability approaches for their implementation.

Impact analysis is used to reason about the effects of a change to a development artefact. As well as facilitating change propagation, impact analysis can help to predict the cost and complexity of changes [Bohner 2002]. Impact analysis requires solutions to several sub-problems, which include change detection, analysis of the effects of a change, and effective display of the analysis.

[Briand *et al.* 2003] contributes an impact analysis tool for UML models that compares original and evolved versions of the same model, producing a report of evolved model elements that have been impacted by the changes to the original model elements. To facilitate the impact analysis, [Briand *et al.* 2003] identifies change patterns that comprise, among other properties, a trigger (for change detection) and an impact rule (for marking model elements affected by this change). Figure 3.4 shows a sample impact analysis pattern for UML sequence diagrams, which is triggered when a message is added, and marks the sending class, the sending operation and the postcondition of the sending operation as impacted.

[Winkler & Pilgrim 2009] note that only event-based approaches, such as the one described in [Briand *et al.* 2003], have been proposed for automating impact analysis. Because of the use of patterns for detecting changes and determining reactions, event-based impact analysis is similar to differencing approaches for schema evolution (for example, [Lerner 2000], which was discussed in Section 3.2.2). When more than one trigger might apply, event-based impact analysis approaches must provide mechanisms for selecting between applicable patterns. In [Briand *et al.* 2003], the selection policy is implicit (cannot be changed by the user) and further analysis is needed to assess its limitations.

Finally, model synchronisation tools might apply techniques used in automated synchronisation tools for traditional development environments, such as the refactoring functionality of the Eclipse Java Development Tools [Fuhrer *et al.* 2007].

**Synchronisation of models with text and trace links** So far, this section has concentrated on model-to-model synchronisation, which is facilitated by traceability. Traceability is important for other software evolution activities in a model-driven development environment – such as synchronisation between models and text and between models and trace links – and these activities are now discussed.

While most of the model synchronisation literature focuses on synchronising models with other models, some papers consider synchronisation between models and other types of artefact. For synchronising changes in requirements documents with models, there is abundance of work in the field of requirements engineering, where the need for traceability was first reported. For synchronising models with generated text (during code generation, for example), the model-to-text language, Epsilon Generation Language (EGL) [Rose *et al.* 2008], produces traceability links between code generation templates and generated files. Sections of code can be marked protected, and are not overwritten by subsequent invocations of the code generation template. As described in [Olsen & Oldevik 2007],

**Change Title:**  Changed Sequence Diagram – Added Message

**Change Code:** CSDVAM

**Changed Element:**  `model::behaviouralElements::collaborations::`
`SequenceDiagramView`

**Added Property:**  `model::behaviouralElements::collaborations::Message`

**Impacted Elements:** `model::foundation::core::ClassClassifier`
`model::foundation::core::Operation`
`model::foundation::core::Postcondition`

**Description:** The base class of the classifier role that sends the added message is impacted. The operation that sends the added message is impacted and its postcondition is also impacted.

**Rationale:** The sending/source class now sends a new message and one of its operations, actually sending the added message, is impacted. This operation is known or not, depending on whether the message triggering the added message corresponds to an invoked operation. If, for example, it is a signal then we may not know the operation, just by looking at the sequence diagram. The impacted postcondition may now not represent the effect (what is true on completion) of its operation.

**Resulting Changes:** The implementation of the base class may have to be modified. The method of the impacted operation may have to be modified. The impacted postcondition should be checked to ensure that it is still valid.

**Invoked Rule:** Changed Class Operation – Changed Postcondition (CCOCPst)

**OCL Expressions:**
```
context modelChanges::Change def:
    let addedMessage:Message = self.changedElement.oclAsType(SequenceDiagramView).
        Message->select(m:Message | m.getIDStr()=self.propertyID)
    let sendingOperation:Operation = (
        if addedMessage.activator.action.oclIsTypeOf(CallAction) then
            addedMessage.sender.base.operation->select(o:Operation |
            o.equals(addedMessage.activator.callAction.operation))
        else
            null
        endif)
context modelChanges::Change — class
    addedMessage.sender.base
context modelChanges::Change — operation
    sendingOperation
context modelChanges::Change — postcondition
    sendingOperation.postcondition
```

Figure 3.4: Exemplar impact analysis pattern, taken from [Briand *et al.* 2003].

the MOFScript model-to-text language, like EGL, provides protected sections and, unlike EGL, also stores traceability links in a structured manner. The traceability links described in [Olsen & Oldevik 2007] can be used for impact analysis, model coverage (for highlighting which areas of the model contribute to the generated code) and orphan analysis (for detecting invalid traceability links).

Trace links can be affected when development artefacts change. Synchronisation tools rely on accurate trace links and hence the maintenance of trace links is important. [Winkler & Pilgrim 2009] suggests that trace versioning should be used to address the challenges of trace link maintenance, which include the accidental inclusion of unintended dependencies as well as the exclusion of necessary dependencies. Furthermore, [Winkler & Pilgrim 2009] notes that, although versioning traces has been explored in specialised areas (such as hypermedia [Nguyen *et al.* 2005]), there is no holistic approach for versioning traces.

**Model-metamodel Co-Evolution**

A metamodel describes the structures and rules for a family of models. When a model uses the structures and adheres to the rules defined by a metamodel, the model is said to *conform* to the metamodel [Bézivin 2005]. A change to a metamodel might require changes to models to ensure the preservation of conformance. The process of evolving a metamodel and its models together to preserve conformance is termed *model-metamodel co-evolution* and is subsequently referred to as *co-evolution*. This section explores existing approaches to co-evolution, comparing them with work from the closely related areas of schema and grammar evolution approaches (Sections 3.2.2 and 3.2.3). A more thorough analysis of co-evolution approaches is conducted in Chapter 4.

**Co-evolution theory**   A co-evolution process involves changing a metamodel and updating instance models to preserve conformance. Often, the two activities are considered separately, and the latter is termed *migration*. In this thesis, the term *migration strategy* is used to mean an algorithm that specifies migration. [Sprinkle & Karsai 2004] were the first to identify the need for approaches that consider the specific requirements of co-evolution, treating it separately from other development artefacts. In particular, Sprinkle and Karsai describe migration as distinct from – and as having unique challenges compared to – the more general activity of model-to-model transformation. [Sprinkle 2003] uses the phrase "evolution, not revolution" to highlight and emphasise that, during co-evolution, the difference between source and target metamodels is often small.

Understanding the situations in which co-evolution must be managed is important for formulating the requirements for co-evolution tools. However, co-evolution literature rarely reports on the ways in which co-evolution is managed in practice. [Herrmannsdoerfer *et al.* 2009b] reports that migration is sometimes made unnecessary by evolving a metamodel such that the conformance of models is not affected (for example, making only additive changes). [Cicchetti *et al.* 2008] suggests that co-evolution can be carried out by more than one person, and that metamodel developers and model users might not know one another.

**Co-evolution patterns**   Much of the co-evolution literature suggests that the way in which migration is performed should vary depending on the type of meta-model changes made [Gruschko *et al.* 2007, Herrmannsdoerfer *et al.* 2009b, Cicchetti *et al.* 2008, Garcés *et al.* 2009]. In particular, the co-evolution literature identifies two important classifications of metamodel changes that affect the way in which migration is performed. [Gruschko *et al.* 2007] classify metamodel changes, recognising that, depending on the type of metamodel change, migration might be unnecessary (*non-breaking* change), can be automated (*breaking and resolvable* change) and can be automated only when guided by a developer (*breaking and non-resolvable* change). [Herrmannsdoerfer *et al.* 2008] classify metamodel changes into *metamodel-independent* (observed in the evolution of more than one metamodel) and *metamodel-specific* (observed in the evolution of only one metamodel).

Further research is needed to identify categories of metamodel changes because automated co-evolution approaches are built atop them. [Herrmannsdoerfer *et al.* 2008] suggests that a large fraction of metamodel changes re-occur, but the study considers only two metamodels, both taken from the same organisation. Assessing the extent to which changes re-occur across a larger and broader range of metamodels is an open research challenge to which this thesis contributes, particularly in Chapter 4.

**Co-evolution approaches**   Several approaches for managing co-evolution have been proposed, most of which are based on one of the two classifications of meta-model changes described above.

Re-use of migration knowledge is a primary concern in the work of Herrmannsdöerfer. [Herrmannsdoerfer *et al.* 2008] describes an empirical study of the history of two metamodels from the automobile industry, observing that a large number of metamodel changes re-occur. [Herrmannsdoerfer *et al.* 2009b] proposes a co-evolution tool, COPE, that provides a library of co-evolutionary operators. Operators are applied to evolve a metamodel and have pre-defined migration semantics. The application of each operator is recorded, and used to generate an executable migration strategy. Due to its use of re-usable operators, COPE shares characteristics with operator-based approaches for schema and grammar evolution (Sections 3.2.2 and 3.2.3). Consequently, the limitations for operator-based schema evolution approaches identified in [Lerner 2000] apply to COPE. Balancing expressiveness and understandability is a key challenge for operator-based approaches because the former implies a large number of operators while the latter a small number of operators.

[Gruschko *et al.* 2007] suggest inferring co-evolution strategies, based on either a difference model of two versions of the evolving metamodel (direct comparison) or on a list of changes recorded during the evolution of a metamodel (indirect comparison). To this end, [Gruschko *et al.* 2007] contributes a *metamodel matching* co-evolution process, shown in Figure 3.5.

Both [Cicchetti *et al.* 2008] and [Garcés *et al.* 2009] extend the work of [Gruschko *et al.* 2007], using the metamodel matching co-evolution process. Both also use higher-order model transformation[3] for determining the migration strategy (the penultimate phase in Figure 3.5). [Cicchetti *et al.* 2008] contributes a metamodel

---

[3]A model-to-model transformation that consumes or produces a model-to-model transformation is *higher-order*.
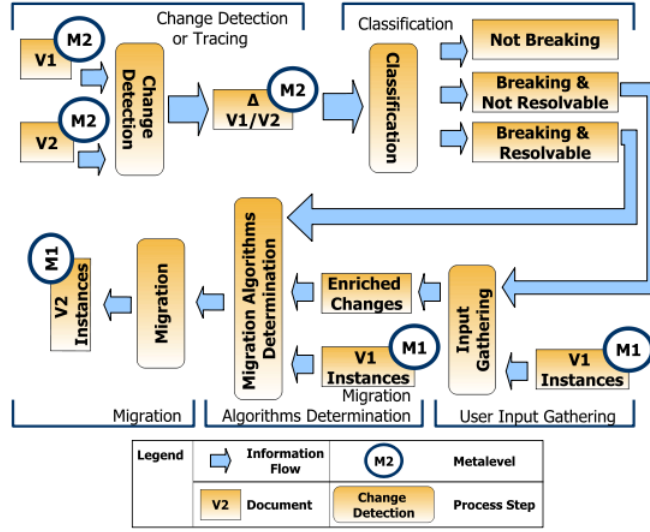
Figure 3.5: Co-evolution process, taken from [Gruschko *et al.* 2007].

for describing the similarities and differences between two versions of a meta-model, enabling a model-driven approach to generating model migration strategies. [Garcés *et al.* 2009] provides a similar metamodel, but uses a metamodel matching process that can be customised by the user, who specifies matching heuristics to form a matching strategy. Otherwise, the co-evolution approaches described in [Cicchetti *et al.* 2008] and [Garcés *et al.* 2009] are fully automatic and cannot be guided by the user. Clearly then, accuracy is important for approaches that compare two metamodel versions, but the co-evolution literature does not assess the extent to which approaches like [Cicchetti *et al.* 2008] and [Garcés *et al.* 2009] can be applied.

Some co-evolution approaches predate the classifications of metamodel changes described above. For instance, [Wachsmuth 2007] proposes a preliminary catalogue of metamodel changes and was the first to employ higher-order transformation for specifying model migration. However, [Wachsmuth 2007] considers a small number of metamodel changes occurring in isolation and, as such, it is not clear whether the approach can be used in general. [Sprinkle 2003] proposes a visual transformation language for specifying model migration, based on graph transformation theory. As such, the migration language proposed by Sprinkle is less expressive than imperative or hybrid transformation languages (as discussed in Section **??**).

**Summary**    Automated migration is still an open research challenge. Co-evolution approaches are in their infancy, and key problems need to be addressed. For example, [Lerner 2000] notes that matching schemas (metamodels) can yield more than one feasible set of migration strategies. [Cicchetti *et al.* 2008] does not acknowledge this challenge. [Garcés *et al.* 2009] offers heuristics for controlling metamodel matching, which might affect the predictability of the co-evolution process.

Another open research challenge is in identifying an appropriate notation for describing migration. [Wachsmuth 2007, Cicchetti *et al.* 2008] use higher-order transformations, while [Herrmannsdoerfer *et al.* 2009b] uses a general-purpose programming language. Because migration is a specialisation of model-to-model transformation [Sprinkle & Karsai 2004], languages other than model-to-model transformation languages might be more suitable for describing migration.

Until co-evolution tools reach maturity, improving MDE modelling frameworks to better support co-evolution is necessary. For example, the Eclipse Modelling Framework [Steinberg *et al.* 2008] cannot load models that no longer conform to their metamodel and, hence non-conformant models cannot be used for model-driven development with EMF.

**Visualisation**

To better understand the effects of evolution on development artefacts, visualising different versions of each artefact may be beneficial. Existing research for comparing text can be enhanced to perform semantic-differencing of models with a textual concrete syntax. For models with a visual concrete syntax, another approach is required.

[Pilgrim *et al.* 2008] have implemented a three-dimensional editor for exploring transformation chains (the sequential composition of model-to-model transformations). Their tool enables developers to visualise the way in which model elements are transformed throughout the chain. Figure 3.6 depicts a sample transformation chain visualisation. Each plane represents a model. The links between each plane illustrates the effects of a model-to-model transformation.

The visualisation technology described in [Pilgrim *et al.* 2008] could be used to facilitate exploration of artefact evolution.

## 3.3   Summary

This chapter reviews and analyses software evolution literature, introducing terminology and describing *impact analysis* and *change propagation*, two evolution activities explored in the remainder of this thesis. Principles and practices of software evolution (from the fields of programming languages, relational database systems and grammarware) were compared, contrasted and analysed. In particular, software evolution literature from the MDE community was reviewed and analysed to allow the formulation of potential research directions for this thesis. The chapter now concludes by synthesising, from the reviewed literature, research challenges for identifying and managing software evolution in the context of MDE.

**Model Refactoring Challenges**   The model refactoring literature propose tools and techniques for improving the quality of existing models without affecting their functional behaviour. In traditional development environments, inter-artefact refactoring (in which changes span more than one development artefact) is often automated, but none of the model refactoring papers discussed in this chapter consider inter-model refactoring. In general, the refactoring literature covers several concerns, such as identification, validation and assessment (Section 3.1.2), but the model refactoring literature considers only the specification

Figure 3.6: Visualising a transformation chain [Pilgrim *et al.* 2008].

and application of refactoring. To better understand the costs and benefits of model refactoring, further model refactoring research must consider all of the concerns considered in the refactoring literature in general.

**Model Synchronisation Challenges**  Improved scalability is the primary motivation of most model synchronisation research. However, [Fritzsche *et al.* 2008] suggest that model synchronisation can be used to improve the maintainability of a system via modularisation. Building on the work by [Fritzsche *et al.* 2008], further research should explore the extent to which model synchronisation can be used to manage evolution. [Winkler & Pilgrim 2009] observe that, for impact analysis between models, only event-based approaches have been reported; other approaches – used successfully to manage evolution in other fields (such as relational databases and grammarware) – have not been applied. Few papers consider synchronisation with other artefacts and maintaining trace links and there is potential for further research in these areas.

**Model-Metamodel Co-evolution Challenges**  To better understand model-metamodel co-evolution, further studies of the ways in which metamodels change are required. [Herrmannsdoerfer *et al.* 2008] report an empirical study of industrial metamodels, but focus only on two metamodels produced in the same organisation.  Challenges for co-evolution reported in other fields have not been addressed by the model-metamodel co-evolution literature. For example, [Lerner 2000] notes that comparing two versions of a changed artefact (such as

metamodel) can suggest more than one feasible migration strategy. Approaches to co-evolution that do not consider the way in which a metamodel has changed, such as [Cicchetti *et al.* 2008, Garcés *et al.* 2009] must address this challenge. A range of notations are used for model migration, including model-to-model transformation languages and general-purpose programming languages, which is a challenge for the comparison of co-evolution tools. Finally, contemporary MDE modelling frameworks do not facilitate MDE for non-conformant models, which is problematic at least until co-evolution tools reach maturity.

**General Challenges for Evolution in MDE**  From the analysis in this chapter, several research challenges for software evolution in the context of MDE are apparent. Greater understanding of the situations in which evolution occurs informs the identification and management of evolution, yet few papers study evolution in real-world MDE projects. Analysis of existing projects can yield patterns of evolution, providing a common vocabulary for thinking and communicating about evolution. Evolution notations and tools are built atop these patterns to automate some evolution activities. Few papers that consider evolution focus less on evolution patterns than on evolution tools and their implementation. Finally, recording, analysing and visualising changes made over the long term to MDE development artefacts and to MDE projects is an area that is not considered in the literature.

As well as directing the thesis research, the above challenges influenced the choice of research method. Most of the software evolution research discussed in this chapter uses a similar method: first, identify and categorise evolutionary changes by considering all of the ways in which artefacts can change. Next, design a taxonomy of operators that capture these changes or a matching algorithm that detects the application of the changes. Then, implement a tool for applying operators, invoking a matching algorithm, or trigger change events. Finally, evaluate the tool on existing projects containing examples of evolution.

The research in this thesis follows a different method, based on the method used by Digg in his work on program refactoring ([Dig & Johnson 2006b, Dig & Johnson 2006a, Dig *et al.* 2006, Dig *et al.* 2007]). First, existing projects are analysed to better understand the situations in which evolution occurs. From this analysis, research requirements are derived, and structures and process for identifying and managing evolution are implemented. The structures and process are evaluated by comparison with related work and by application on an existing project in which there is a need to identify and manage evolution.

Using the literature reviewed and the research challenges identified in this chapter, Chapter 4 analyses examples of evolution from existing MDE projects and derives requirements for structures and processes for identifying and managing evolution in the context of MDE.

# Bibliography

[37-Signals 2008]   37-Signals. Ruby on Rails [online]. [Accessed 30 June 2008] Available at: `http://www.rubyonrails.org/`, 2008.

[Aizenbud-Reshef *et al.* 2005]   N. Aizenbud-Reshef, R.F. Paige, J. Rubin, Y. Shaham-Gafni, and D.S. Kolovos. Operational semantics for traceability. In *Proc. ECMDA-FA Workshop on Traceability*, pages 8–14, 2005.

[Alexander *et al.* 1977]   Christopher Alexander, Sara Ishikawa, and Murray Silverstein. *A Pattern Language: Towns, Buildings, Construction (Center for Environmental Structure Series)*. Oxford University Press, 1977.

[Álvarez *et al.* 2001]   José Álvarez, Andy Evans, and Paul Sammut. MML and the metamodel architecture. In *Proc. Workshop on Transformation in UML*, 2001.

[Arendt *et al.* 2009]   Thorsten Arendt, Florian Mantz, Lars Schneider, and Gabriele Taentzer. Model refactoring in eclipse by LTK, EWL, and EMF Refactor: A case study. In *Proc. Joint MoDSE-MCCM Workshop*, 2009.

[ATLAS 2007]   ATLAS. Atlas Transformation Language Project Website [online]. [Accessed 30 June 2008] Available at: `http://www.eclipse.org/m2m/atl/`, 2007.

[Backus 1978]   John Backus. The history of FORTRAN I, II and III. *History of Programming Languages*, 1:165–180, 1978.

[Balazinska *et al.* 2000]   Magdalena Balazinska, Ettore Merlo, Michel Dagenais, Bruno Lagüe, and Kostas Kontogiannis. Advanced clone-analysis to support object-oriented system refactoring. In *Proc. Working Conference on Reverse Engineering*, pages 98–107. IEEE Computer Society, 2000.

[Banerjee *et al.* 1987]   Jay Banerjee, Won Kim, Hyoung-Joo Kim, and Henry F. Korth. Semantics and implementation of schema evolution in object-oriented databases. In *Proc. Special Interest Group on Management of Data*, volume 16, pages 311–322. ACM, 1987.

[Beck & Cunningham 1989]   Kent Beck and Ward Cunningham. Constructing abstractions for object-oriented applications. *Journal of Object Oriented Programming*, 2, 1989.

[Bézivin 2005]   Jean Bézivin. On the unification power of models. *Software and System Modeling*, 4(2):171–188, 2005.

[Biermann *et al.* 2006]     Enrico Biermann, Karsten Ehrig, Christian Köhler, Günter Kuhns, Gabriele Taentzer, and Eduard Weiss. Emf model refactoring based on graph transformation concepts. *ECEASST*, 3, 2006.

[Bloch 2005]     Joshua Bloch. How to design a good API and why it matters [online]. Keynote address to the LCSD Workshop at OOPSLA, October 2005, San Diego, United States of America. [Accessed 23 July 2009] Available at: `http://lcsd05.cs.tamu.edu/slides/keynote.pdf`, 2005.

[Bohner 2002]     Shawn A. Bohner. Software change impacts - an evolving perspective. In *Proc. International Conference on Software Maintenance (ICSM)*, pages 263–272. IEEE Computer Society, 2002.

[Bosch 1998]     Jan Bosch. Design patterns as language constructs. *Journal of Object Oriented Programming*, 11(2):18–32, 1998.

[Briand *et al.* 2003]     Lionel C. Briand, Yvan Labiche, and L. O'Sullivan. Impact analysis and change management of uml models. In *Proc. International Conference on Software Maintenance (ICSM)*, pages 256–265. IEEE Computer Society, 2003.

[Brown *et al.* 1998]     William J. Brown, Raphael C. Malveau, Hays W. McCormick III, and Thomas J. Mowbray. *Anti Patterns*. Wiley, 1998.

[Cervelle *et al.* 2006]     Julien Cervelle, Rémi Forax, and Gilles Roussel. Tatoo: an innovative parser generator. In *Principles and Practice of Programming in Java*, pages 13–20. ACM, 2006.

[Ceteva 2008]     Ceteva. XMF – the extensible programming language [online]. [Accessed 30 June 2008] Available at: `http://www.ceteva.com/xmf.html`, 2008.

[Chen & Chou 1999]     J.Y.J. Chen and S.C. Chou. Consistency management in a process environment. *Systems and Software*, 47(2-3):105–110, 1999.

[Cicchetti *et al.* 2008]     Antonio Cicchetti, Davide Di Ruscio, Romina Eramo, and Alfonso Pierantonio. Automating co-evolution in model-driven engineering. In *Proc. EDOC*, pages 222–231. IEEE Computer Society, 2008.

[Clark *et al.* 2008]     Tony Clark, Paul Sammut, and James Willians. Superlanguages: Developing languages and applications with XMF [online]. [Accessed 30 June 2008] Available at: `http://www.ceteva.com/docs/Superlanguages.pdf`, 2008.

[Cleland-Huang *et al.* 2003]     Jane Cleland-Huang, Carl K. Chang, and Mark Christensen. Event-based traceability for managing evolutionary change. *IEEE Transactions on Software Engineering*, 29(9):796–810, 2003.

[Costa & Silva 2007]     M. Costa and A.R. da Silva. RT-MDD framework – a practical approach. In *Proc. ECMDA-FA Workshop on Traceability*, pages 17–26, 2007.

[Czarnecki & Helsen 2006]     Krzysztof Czarnecki and Simon Helsen. Feature-based survey of model transformation approaches. *IBM Systems Journal*, 45(3):621–646, 2006.

[Deursen *et al.* 2000]    Arie van Deursen, Paul Klint, and Joost Visser. Domain-specific languages: An annotated bibliography. *SIGPLAN Notices*, 35(6):26–36, 2000.

[Deursen *et al.* 2007]    Arie van Deursen, Eelco Visser, and Jos Warmer. Model-driven software evolution: A research agenda. In *Proc. Workshop on Model-Driven Software Evolution*, pages 41–49, 2007.

[Dig & Johnson 2006a]    Danny Dig and Ralph Johnson. Automated upgrading of component-based applications. In *OOPSLA Companion*, pages 675–676, 2006.

[Dig & Johnson 2006b]    Danny Dig and Ralph Johnson. How do APIs evolve? A story of refactoring. *Journal of Software Maintenance and Evolution*, 18(2):83–107, 2006.

[Dig *et al.* 2006]    Danny Dig, Can Comertoglu, Darko Marinov, and Ralph Johnson. Automated detection of refactorings in evolving components. In *Proc. European Conference on Object-Oriented Programming*, volume 4067 of *LNCS*, pages 404–428. Springer, 2006.

[Dig *et al.* 2007]    Danny Dig, Kashif Manzoor, Ralph Johnson, and Tien N. Nguyen. Refactoring-aware configuration management for object-oriented programs. In *Proc. International Conference on Software Engineering*, pages 427–436. IEEE Computer Society, 2007.

[Dmitriev 2004]    Sergey Dmitriev. Language oriented programming: The next programming paradigm. *JetBrains onBoard [online]*, 1, 2004. [Accessed 30 June 2008] Available at: `http://www.onboard.jetbrains.com/is1/articles/04/10/lop/`.

[Drivalos *et al.* 2008]    Nicholas Drivalos, Richard F. Paige, Kiran J. Fernandes, and Dimitrios S. Kolovos. Towards rigorously defined model-to-model traceability. In *Proc. European Conference on the Model Driven Architecture Workshop on Traceability*, 2008.

[Ducasse *et al.* 1999]    Stéphane Ducasse, Matthias Rieger, and Serge Demeyer. A language independent approach for detecting duplicated code. In *Proc. International Conference on Software Maintenance*, pages 109–118. IEEE Computer Society, 1999.

[Eclipse 2008a]    Eclipse. Eclipse Modeling Framework project [online]. [Accessed 22 January 2009] Available at: `http://www.eclipse.org/modeling/emf/`, 2008.

[Eclipse 2008b]    Eclipse. Eclipse project [online]. [Accessed 20 January 2009] Available at: `http://www.eclipse.org`, 2008.

[Eclipse 2008c]    Eclipse. Epsilon home page [online]. [Accessed 30 June 2008] Available at: `http://www.eclipse.org/gmt/epsilon/`, 2008.

[Eclipse 2008d]    Eclipse. Generative Modelling Technologies project [online]. [Accessed 30 June 2008] Available at: `http://www.eclipse.org/gmt`, 2008.

[Eclipse 2009a]    Eclipse. Model Development Tools project [online]. [Accessed 6 January 2008] Available at: `http://www.eclipse.org/modeling/mdt/`, 2009.

[Eclipse 2009b]    Eclipse. UML2 Model Development Tools project [online]. [Accessed 7 September 2009] Available at: `http://www.eclipse.org/modeling/mdt/uml2`, 2009.

[Eclipse 2010]    Eclipse. Connected Data Objects Model Repository Project Website [online]. [Accessed 15 February 2010] Available at: `http://www.eclipse.org/modeling/emf/?project=cdo#cdo`, 2010.

[Edelweiss & Freitas Moreira 2005]    Nina Edelweiss and Álvaro Freitas Moreira. Temporal and versioning model for schema evolution in object-oriented databases. *Data & Knowledge Engineering*, 53(2):99–128, 2005.

[Elmasri & Navathe 2006]    Ramez Elmasri and Shamkant B. Navathe. *Fundamentals of Database Systems*. Addison-Wesley Longman, 2006.

[Erlikh 2000]    Len Erlikh. Leveraging legacy system dollars for e-business. *IT Professional*, 2(3):17–23, 2000.

[Evans 2004]    E. Evans. *Domain-Driven Design: Tacking Complexity In the Heart of Software*. Addison-Wesley, Boston, MA, USA, 2004.

[Ferrandina *et al.* 1995]    Fabrizio Ferrandina, Thorsten Meyer, Roberto Zicari, Guy Ferran, and Joëlle Madec. Schema and database evolution in the O2 object database system. In *Very Large Data Bases*, pages 170–181. Morgan Kaufmann, 1995.

[Fowler 1999]    Martin Fowler. *Refactoring: improving the design of existing code*. Addison-Wesley, 1999.

[Fowler 2002]    Martin Fowler. *Patterns of Enterprise Application Architecture*. Addison-Wesley, 2002.

[Fowler 2005]    Martin Fowler. Language workbenches: The killer-app for domain specific languages? [online]. [Accessed 30 June 2008] Available at: `http://www.martinfowler.com/articles/languageWorkbench.html`, 2005.

[Fowler 2010]    Martin Fowler. *Domain Specific Languages*. Addison-Wesley Professional, 2010.

[Frankel 2002]    David Frankel. *Model Driven Architecture: Applying MDA to Enterprise Computing*. Wiley, 2002.

[Fritzsche *et al.* 2008]    M. Fritzsche, J. Johannes, S. Zschaler, A. Zherebtsov, and A. Terekhov. Application of tracing techniques in Model-Driven Performance Engineering. In *Proc. ECMDA Traceability Workshop (ECMDA-TW)*, pages 111–120, 2008.

[Fuhrer *et al.* 2007]    Robert M. Fuhrer, Adam Kiezun, and Markus Keller. Refactoring in the Eclipse JDT: Past, present, and future. In *Proc. Workshop on Refactoring Tools*, 2007.

[Gamma *et al.* 1995]   Erich Gamma, Richard Helm, Ralph Johnson, and John Vlissides. *Design patterns: elements of reusable object-oriented software.* Addison-Wesley, 1995.

[Garcés *et al.* 2009]   Kelly Garcés, Frédéric Jouault, Pierre Cointe, and Jean Bézivin. Managing model adaptation by precise detection of metamodel changes. In *Proc. ECMDA-FA*, volume 5562 of *LNCS*, pages 34–49. Springer, 2009.

[Gosling *et al.* 2005]   James Gosling, Bill Joy, Guy Steele, and Gilad Bracha. *The Java$^{TM}$ Language Specification.* Addison-Wesley, Boston, MA, USA, 2005.

[Graham 1993]   Paul Graham. *On Lisp: Advanced Techniques for Common Lisp.* Prentice-Hall, 1993.

[Greenfield *et al.* 2004]   Jack Greenfield, Keith Short, Steve Cook, and Stuart Kent. *Software Factories: Assembling Applications with Patterns, Models, Frameworks, and Tools.* Wiley, 2004.

[Gronback 2009]   R.C. Gronback. *Eclipse Modeling Project: A Domain-Specific Language (DSL) Toolkit.* Addison-Wesley Professional, 2009.

[Gruschko *et al.* 2007]   Boris Gruschko, Dimitrios S. Kolovos, and Richard F. Paige. Towards synchronizing models with evolving metamodels. In *Proc. Workshop on Model-Driven Software Evolution*, 2007.

[Guerrini *et al.* 2005]   Giovanna Guerrini, Marco Mesiti, and Daniele Rossi. Impact of XML schema evolution on valid documents. In *Proc. Workshop on Web Information and Data Management*, pages 39–44, 2005.

[Hearnden *et al.* 2006]   David Hearnden, Michael Lawley, and Kerry Raymond. Incremental model transformation for the evolution of model-driven systems. In *Model Driven Engineering Languages and Systems*, volume 4199 of *LNCS*, pages 321–335. Springer, 2006.

[Heidenreich *et al.* 2009]   Florian Heidenreich, Jendrik Johannes, Sven Karol, Mirko Seifert, and Christian Wende. Derivation and refinement of textual syntax for models. In *Proc. ECMDA-FA*, volume 5562 of *LNCS*, pages 114–129. Springer, 2009.

[Herrmannsdoerfer *et al.* 2008]   Markus Herrmannsdoerfer, Sebastian Benz, and Elmar Juergens. Automatability of coupled evolution of metamodels and models in practice. In *Proc. International Conference on Model Driven Engineering Languages and Systems*, volume 5301 of *LNCS*, pages 645–659. Springer, 2008.

[Herrmannsdoerfer *et al.* 2009a]   M. Herrmannsdoerfer, D. Ratiu, and G. Wachsmuth. Language evolution in practice. In *Proc. SLE*, volume 5696 of *LNCS*, pages 3–22. Springer, 2009.

[Herrmannsdoerfer *et al.* 2009b]   Markus Herrmannsdoerfer, Sebastian Benz, and Elmar Juergens. COPE - automating coupled evolution of metamodels and models. In *Proc. ECOOP*, volume 5653 of *LNCS*, pages 52–76. Springer, 2009.

[Hussey & Paternostro 2006]    Kenn Hussey and Marcelo Paternostro. Advanced features of EMF. Tutorial at EclipseCon 2006, California, USA. [Accessed 07 September 2009] Available at: `http://www.eclipsecon.org/2006/Sub.do?id=171`, 2006.

[IBM 2005]    IBM. Emfatic Language for EMF Development [online]. [Accessed 30 June 2008] Available at: `http://www.alphaworks.ibm.com/tech/emfatic`, 2005.

[IRISA 2007]    IRISA. Sintaks. `http://www.kermeta.org/sintaks/`, 2007.

[ISO/IEC 1996]    Information Technology ISO/IEC. Syntactic metalanguage – Extended BNF. ISO 14977:1996 International Standard, 1996.

[ISO/IEC 2002]    Information Technology ISO/IEC. Z Formal Specification Notation – Syntax, Type System and Semantics. ISO 13568:2002 International Standard, 2002.

[JetBrains 2008]    JetBrains. MPS – Meta Programming System [online]. [Accessed 30 June 2008] Available at: `http://www.jetbrains.com/mps/index.html`, 2008.

[Jouault & Kurtev 2005]    Frédéric Jouault and Ivan Kurtev. Transforming models with ATL. In *Proc. Satellite Events at the International Conference on Model Driven Engineering Languages and Systems*, volume 3844 of *LNCS*, pages 128–138. Springer, 2005.

[Jouault 2005]    Frédéric Jouault. Loosely coupled traceability for ATL. In *Proc. ECMDA-FA Workshop on Traceability*, 2005.

[Kataoka *et al.* 2001]    Yoshio Kataoka, Michael D. Ernst, William G. Griswold, and David Notkin. Automated support for program refactoring using invariants. In *Proc. International Conference on Software Maintenance*, pages 736–743. IEEE Computer Society, 2001.

[Kelly & Tolvanen 2008]    Steven Kelly and Juha-Pekka Tolvanen. *Domain-Specific Modelling*. Wiley, 2008.

[Kerievsky 2004]    Joshua Kerievsky. *Refactoring to Patterns*. Addison-Wesley, 2004.

[Kleppe *et al.* 2003]    Anneke G. Kleppe, Jos Warmer, and Wim Bast. *MDA Explained: The Model Driven Architecture: Practice and Promise*. Addison-Wesley, 2003.

[Klint *et al.* 2003]    P. Klint, R. Lämmel, and C. Verhoef. Towards an engineering discipline for grammarware. *ACM Transactions on Software Engineering Methodology*, 14:331–380, 2003.

[Kolovos *et al.* 2006a]    Dimitrios S. Kolovos, Richard F. Paige, and Fiona A.C. Polack. Merging models with the epsilon merging language (eml). In *Proc. MoDELS*, volume 4199 of *Lecture Notes in Computer Science*, pages 215–229. Springer, 2006.

[Kolovos *et al.* 2006b]   Dimitrios S. Kolovos, Richard F. Paige, and Fiona A.C. Polack. Model comparison: a foundation for model composition and model transformation testing. In *Proc. Workshop on Global Integrated Model Management*, pages 13–20, 2006.

[Kolovos *et al.* 2006c]   Dimitrios S. Kolovos, Richard F. Paige, and Fiona A.C. Polack. The Epsilon Object Language (EOL). In *Proc. ECMDA-FA*, volume 4066 of *LNCS*, pages 128–142. Springer, 2006.

[Kolovos *et al.* 2007]   Dimitrios S. Kolovos, Richard F. Paige, Fiona A.C. Polack, and Louis M. Rose. Update transformations in the small with the Epsilon Wizard Language. *Journal of Object Technology*, 6(9):53–69, 2007.

[Kolovos *et al.* 2008a]   Dimitrios S. Kolovos, Richard F. Paige, and Fiona Polack. The Epsilon Transformation Language. In *Proc. ICMT*, volume 5063 of *LNCS*, pages 46–60. Springer, 2008.

[Kolovos *et al.* 2008b]   Dimitrios S. Kolovos, Richard F. Paige, and Fiona A.C. Polack. On the evolution of OCL for capturing structural constraints in modelling languages. In *Proc. Workshop on Rigorous Methods for Software Construction and Analysis*, 2008.

[Kolovos *et al.* 2008c]   Dimitrios S. Kolovos, Richard F. Paige, and Fiona A.C. Polack. Scalability : The holy grail of model driven engineering. In *Proc. Workshop on Challenges in Model Driven Engineering*, 2008.

[Kolovos *et al.* 2009]   Dimitrios S. Kolovos, Richard F. Paige, and Louis M. Rose. EuGENia: GMF for mortals. Long talk at Eclipse Summit Europe, October 2009, Ludwigsburg, Germany. Available at: `https://www.eclipsecon.org/submissions/ese2009/view_talk.php?id=979` [Accessed 12 April 2010], 2009.

[Kolovos 2009]   Dimitrios S. Kolovos. *An Extensible Platform for Specification of Integrated Languages for Model Management*. PhD thesis, University of York, United Kingdom, 2009.

[Kramer 2001]   Diane Kramer. XEM: XML Evolution Management. Master's thesis, Worcester Polytechnic Institute, MA, USA, 2001.

[Lago *et al.* 2009]   Patricia Lago, Henry Muccini, and Hans van Vliet. A scoped approach to traceability management. *Systems and Software*, 82(1):168–182, 2009.

[Lämmel & Verhoef 2001]   R. Lämmel and C. Verhoef. Semi-automatic grammar recovery. *Software - Practice and Experience*, 31(15):1395–1438, 2001.

[Lämmel 2001]   R. Lämmel. Grammar adaptation. In *Proc. Formal Methods for Increasing Software Productivity (FME), International Symposium of Formal Methods Europe*, volume 2021 of *LNCS*, pages 550–570. Springer, 2001.

[Lämmel 2002]   R. Lämmel. Towards generic refactoring. In *Proc. ACM SIG-PLAN Workshop on Rule-Based Programming*, pages 15–28. ACM, 2002.

[Lehman 1969]    Meir M. Lehman. The programming process. Technical report, IBM Res. Rep. RC 2722, 1969.

[Lehman 1978]    Meir M. Lehman. Programs, cities, students - limits to growth? *Programming Methodology*, pages 42–62, 1978.

[Lehman 1980]    Meir M. Lehman. On understanding laws, evolution, and conservation in the large-program life cycle. *Journal of Systems and Software*, 1:213–221, 1980.

[Lehman 1985]    Meir M. Lehman. *Program evolution: processes of software change.* Academic, 1985.

[Lehman 1996]    Meir M. Lehman. Laws of software evolution revisited. In *Proc. European Workshop on Software Process Technology*, pages 108–124, 1996.

[Lerner 2000]    Barbara Staudt Lerner. A model for compound type changes encountered in schema evolution. *ACM Transactions on Database Systems*, 25(1):83–127, 2000.

[Mäder *et al.* 2008]    P. Mäder, O. Gotel, and I. Philippow. Rule-based maintenance of post-requirements traceability relations. In *Proc. IEEE International Requirements Engineering Conference (RE)*, pages 23–32, 2008.

[Martin & Martin 2006]    R.C. Martin and M. Martin. *Agile Principles, Patterns, and Practices in C#.* Prentice Hall, Upper Saddle River, NJ, USA, 2006.

[McCarthy 1978]    John McCarthy. History of Lisp. *History of Programming Languages*, 1:217–223, 1978.

[McNeile 2003]    Ashley McNeile. MDA: The vision with the hole? [Accessed 30 June 2008] Available at: `http://www.metamaxim.com/download/documents/MDAv1.pdf`, 2003.

[Mellor & Balcer 2002]    Stephen J. Mellor and Marc Balcer. *Executable UML: A Foundation for Model-Driven Architectures.* Addison-Wesley Longman, 2002.

[Mens & Tourwé 2004]    Tom Mens and Tom Tourwé. A survey of software refactoring. *IEEE Transactions on Software Engineering*, 30(2):126–139, 2004.

[Mens *et al.* 2007]    Tom Mens, Gabriele Taentzer, and Dirk Müller. Challenges in model refactoring. In *Proc. Workshop on Object-Oriented Reengineering*, 2007.

[Merriam-Webster 2010]    Merriam-Webster. Definition of Nuclear Family. `http://www.merriam-webster.com/dictionary/nuclear%20family`, 2010.

[Moad 1990]    J Moad. Maintaining the competitive edge. *Datamation*, 36(4):61–66, 1990.

[Moha *et al.* 2009]   Naouel Moha, Vincent Mahé, Olivier Barais, and Jean-Marc Jézéquel.  Generic model refactorings.  In *Proc. MoDELS*, volume 5795 of *LNCS*, pages 628–643. Springer, 2009.

[Muller & Hassenforder 2005]   Pierre-Alain Muller and Michel Hassenforder. HUTN as a Bridge between ModelWare and GrammarWare.  In *Proc. Workshop in Software Modelling Engineering*, 2005.

[Nentwich *et al.* 2003]   C. Nentwich, W. Emmerich, A. Finkelstein, and E. Ellmer. Flexible consistency checking. *ACM Transactions on Software Engineering and Methodology*, 12(1):28–63, 2003.

[Nguyen *et al.* 2005]   Tien Nhut Nguyen, Cheng Thao, and Ethan V. Munson. On product versioning for hypertexts. In *Proc. International Workshop on Software Configuration Management (SCM)*, pages 113–132. ACM, 2005.

[Oldevik *et al.* 2005]   Jon Oldevik, Tor Neple, Roy Grønmo, Jan Øyvind Aagedal, and Arne-Jørgen Berre. Toward standardised model to text transformations.  In *Proc. ECMDA-FA*, volume 3748 of *LNCS*, pages 239–253. Springer, 2005.

[Olsen & Oldevik 2007]   Gøran K. Olsen and Jon Oldevik. Scenarios of traceability in model to text transformations.  In *Proc. ECMDA-FA*, volume 4530 of *Lecture Notes in Computer Science*, pages 144–156. Springer, 2007.

[OMG 2001]   OMG. Unified Modelling Language 1.4 Specification [online]. [Accessed 15 September 2008] Available at: `http://www.omg.org/spec/UML/1.4/`, 2001.

[OMG 2004]   OMG. Human-Usable Textual Notation 1.0 Specification [online].  [Accessed 30 June 2008] Available at: `http://www.omg.org/technology/documents/formal/hutn.htm`, 2004.

[OMG 2005]   OMG. MOF QVT Final Adopted Specication [online]. [Accessed 22 July 2009] Available at: `www.omg.org/docs/ptc/05-11-01.pdf`, 2005.

[OMG 2006]   OMG. Object Constraint Language 2.0 Specification [online].  [Accessed 30 June 2008] Available at: `http://www.omg.org/technology/documents/formal/ocl.htm`, 2006.

[OMG 2007a]   OMG. Unified Modelling Language 2.1.2 Specification [online]. [Accessed 30 June 2008] Available at: `http://www.omg.org/spec/UML/2.1.2/`, 2007.

[OMG 2007b]   OMG. Unified Modelling Language 2.2 Specification [online]. [Accessed 5 March 2010] Available at: `http://www.omg.org/spec/UML/2.2/`, 2007.

[OMG 2007c]   OMG. XML Metadata Interchange 2.1.1 Specification [online].  [Accessed 30 June 2008] Available at: `http://www.omg.org/technology/documents/formal/xmi.htm`, 2007.

[OMG 2008a]   OMG. Meta-Object Facility [online]. [Accessed 30 June 2008] Available at: `http://www.omg.org/mof`, 2008.

[OMG 2008b]    OMG. Model Driven Architecture [online]. [Accessed 30 June 2008] Available at: `http://www.omg.org/mda/`, 2008.

[OMG 2008c]    OMG. Object Management Group home page [online]. [Accessed 30 June 2008] Available at: `http://www.omg.org`, 2008.

[Opdyke 1992]    William F. Opdyke. *Refactoring Object-Oriented Frameworks.* PhD thesis, University of Illinois at Urbana-Champaign, IL, USA, 1992.

[openArchitectureWare 2007]    openArchitectureWare. openArchitectureWare Project Website [online]. [Accessed 30 June 2008] Available at: `http://www.eclipse.org/gmt/oaw/`, 2007.

[Paige *et al.* 2009]    Richard F. Paige, Louis M. Rose, Xiaocheng Ge, Dimitrios S. Kolovos, and Phillip J. Brooke. FPTC: Automated safety analysis for domain-specific languages. In *MoDELS Workshops and Symposia*, volume 5421 of *Lecture Notes in Computer Science*, pages 229–242. Springer, 2009.

[Parr 2007]    Terence Parr. *The Definitive ANTLR Reference: Building Domain-Specific Languages.* Pragmatic Programmers, 2007.

[Pilgrim *et al.* 2008]    Jens von Pilgrim, Bert Vanhooff, Immo Schulz-Gerlach, and Yolande Berbers. Constructing and visualizing transformation chains. In *Proc. European Conference on the Model Driven Architecture – Foundations and Applications*, volume 5095 of *LNCS*, pages 17–32. Springer, 2008.

[Pizka & Jürgens 2007]    M. Pizka and E. Jürgens. Automating language evolution. In *Proc. Joint IEEE/IFIP Symposium on Theoretical Aspects of Software Engineering (TASE)*, pages 305–315. IEEE Computer Society, 2007.

[Pool 1997]    R. Pool. *Beyond Engineering: How Society Shapes Technology.* Oxford University Press, 1997.

[Porres 2003]    Ivan Porres. Model refactorings as rule-based update transformations. In *Proc. UML*, volume 2863 of *LNCS*, pages 159–174. Springer, 2003.

[Ráth *et al.* 2008]    István Ráth, Gábor Bergmann, András Ökrös, and Dániel Varró. Live model transformations driven by incremental pattern matching. In *Proc. ICMT*, volume 5063 of *LNCS*, pages 107–121. Springer, 2008.

[Rising 2001]    Linda Rising, editor. *Design patterns in communications software.* Cambridge University Press, 2001.

[Rose *et al.* 2008]    Louis M. Rose, Richard F. Paige, Dimitrios S. Kolovos, and Fiona A.C. Polack. The Epsilon Generation Language. In *Proc. European Conference on Model Driven Architecture – Foundations and Applications*, volume 5095 of *LNCS*, pages 1–16. Springer, 2008.

[Rose *et al.* 2009a]    Louis M. Rose, Dimitrios S. Kolovos, Richard F. Paige, and Fiona A.C. Polack. Enhanced automation for managing model and metamodel inconsistency. In *Proc. ASE*. ACM Press, 2009.

[Rose *et al.* 2009b]    Louis M. Rose, Richard F. Paige, Dimitrios S. Kolovos, and Fiona A.C. Polack. An analysis of approaches to model migration. In *Proc. Joint MoDSE-MCCM Workshop*, 2009.

[Rose *et al.* 2010a]    Louis M. Rose, Markus Herrmannsdoerfer, James R. Williams, Dimitrios S. Kolovos, Kelly Garcés, Richard F. Paige, and Fiona A.C. Polack. A comparison of model migration tools. In *Proc. MoDELS*, volume TBC of *Lecture Notes in Computer Science*, page TBC. Springer, 2010.

[Rose *et al.* 2010b]    Louis M. Rose, Dimitrios S. Kolovos, Nicholas Drivalos, James. R. Williams, Richard F. Paige, Fiona A.C. Polack, and Kiran J. Fernandes. Concordance: An efficient framework for managing model integrity [submitted to]. In *Proc. European Conference on Modelling Foundations and Applications*, 2010.

[Rose *et al.* 2010c]    Louis M. Rose, Dimitrios S. Kolovos, Richard F. Paige, and Fiona A.C. Polack. Migrating activity diagrams with epsilon flock. In *Proc. TTC*, 2010.

[Rose *et al.* 2010d]    Louis M. Rose, Dimitrios S. Kolovos, Richard F. Paige, and Fiona A.C. Polack. Model migration case. In *Proc. TTC*, 2010.

[Rose *et al.* 2010e]    Louis M. Rose, Dimitrios S. Kolovos, Richard F. Paige, and Fiona A.C. Polack. Model migration with epsilon flock. In *Proc. ICMT*, volume 6142 of *Lecture Notes in Computer Science*, pages 184–198. Springer, 2010.

[Selic 2003]    Bran Selic. The pragmatics of Model-Driven Development. *IEEE Software*, 20(5):19–25, 2003.

[Selic 2005]    Bran Selic. Whats new in UML 2.0? *IBM Rational software*, 2005.

[Sendall & Kozaczynski 2003]    Shane Sendall and Wojtek Kozaczynski. Model transformation: The heart and soul of model-driven software development. *IEEE Software*, 20:42–45, 2003.

[Sjøberg 1993]    Dag I.K. Sjøberg. Quantifying schema evolution. *Information & Software Technology*, 35(1):35–44, 1993.

[Sommerville 2006]    Ian Sommerville. *Software Engineering*. Addison-Wesley Longman, 2006.

[Sprinkle & Karsai 2004]    Jonathan Sprinkle and Gábor Karsai. A domain-specific visual language for domain model evolution. *Journal of Visual Languages and Computing*, 15(3-4):291–307, 2004.

[Sprinkle 2003]    Jonathan Sprinkle. *Metamodel Driven Model Migration*. PhD thesis, Vanderbilt University, TN, USA, 2003.

[Sprinkle 2008]    Jonathan Sprinkle. *Difference Representation and Conflict Management in Model-Driven Engineering*. PhD thesis, Universita' degli Studi dell'Aquila, L'Aquila, Italy, 2008.

[Stahl *et al.* 2006]    Thomas Stahl, Markus Voelter, and Krzysztof Czarnecki. *Model-Driven Software Development: Technology, Engineering, Management.* Wiley, 2006.

[Steinberg *et al.* 2008]    Dave Steinberg, Frank Budinsky, Marcelo Paternostro, and Ed Merks. *EMF: Eclipse Modeling Framework.* Addison-Wesley Professional, 2008.

[Su *et al.* 2001]    Hong Su, Diane Kramer, Li Chen, Kajal T. Claypool, and Elke A. Rundensteiner. XEM: Managing the evolution of XML documents. In *Proc. Workshop on Research Issues in Data Engineering*, pages 103–110, 2001.

[Tratt 2008]    Laurence Tratt. A change propagating model transformation language. *Journal of Object Technology*, 7(3):107–124, 2008.

[Varró & Balogh 2007]    Dániel Varró and András Balogh. The model transformation language of the VIATRA2 framework. *Science of Computer Programming*, 68(3):187–207, 2007.

[Vries & Roddick 2004]    Denise de Vries and John F. Roddick. Facilitating database attribute domain evolution using meso-data. In *Proc. Workshop on Evolution and Change in Data Management*, pages 429–440, 2004.

[W3C 2007a]    W3C. W3C XML Schema 1.1 Specification [online]. [Accessed 30 June 2008] Available at: `http://www.w3.org/XML/Schema`, 2007.

[W3C 2007b]    W3C. World Wide Web Consortium [online]. [Accessed 30 June 2008] Available at: `http://www.w3.org/`, 2007.

[Wachsmuth 2007]    Guido Wachsmuth. Metamodel adaptation and model co-adaptation. In *Proc. ECOOP*, volume 4609 of *LNCS*, pages 600–624. Springer, 2007.

[Wallace 2005]    Malcolm Wallace. Modular architectural representation and analysis of fault propagation and transformation. *Electronic Notes in Theoretical Computer Science*, 141(3):53–71, 2005.

[Ward 1994]    Martin P. Ward. Language-oriented programming. *Software — Concepts and Tools*, 15(4):147–161, 1994.

[Watson 2008]    Andrew Watson. A brief history of MDA. *Upgrade*, 9(2), 2008.

[Winkler & Pilgrim 2009]    Stefan Winkler and Jens von Pilgrim. A survey of traceability in requirements engineering and model-driven development. *Software and Systems Modeling*, December 2009.